

## clustering\_iris

November 11, 2025

```
[173]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[174]: #scaling, normalization
from sklearn.preprocessing import StandardScaler, MinMaxScaler

#kmeans, dbSCAN, hierarchical (sklearn)
from sklearn.cluster import KMeans, DBSCAN, AgglomerativeClustering
#evaluation
from sklearn.metrics import silhouette_score
#import dataset
from sklearn.datasets import load_iris

#distance matrix (dbSCAN elbow, hierarchical)
from scipy.spatial.distance import pdist, squareform

# hierarchical (scipy)
from scipy.cluster.hierarchy import linkage, dendrogram, fcluster
```

```
[175]: frame = load_iris(as_frame=True)
df = frame['data']
X = df.values
y = np.array(frame['target'])

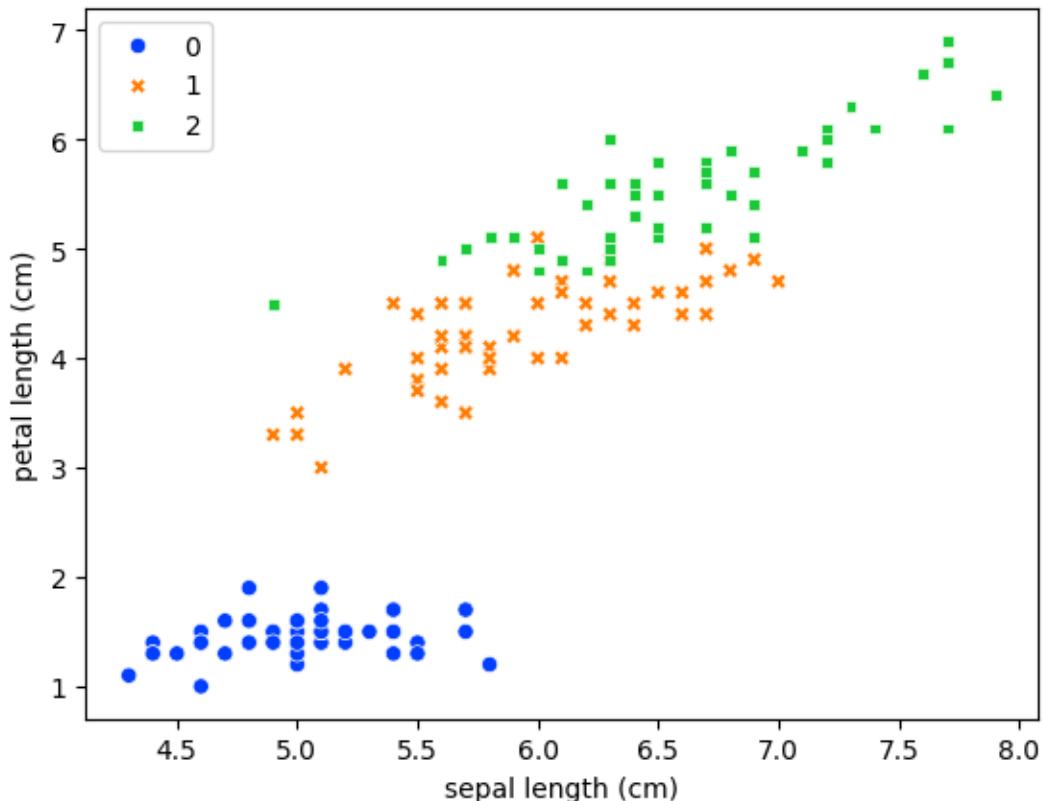
y_map = {0: "setosa", 1: "versicolor", 2: "virginica"}
y_mapped = pd.DataFrame(y).iloc[:,0].map(y_map)

df.head()
```

```
[175]:   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
0              5.1          3.5            1.4            0.2
1              4.9          3.0            1.4            0.2
2              4.7          3.2            1.3            0.2
3              4.6          3.1            1.5            0.2
4              5.0          3.6            1.4            0.2
```

```
[175]:
```

```
[176]: sns.scatterplot(data=df,
                      x="sepal length (cm)",
                      y="petal length (cm)",
                      hue=y,
                      style=y,
                      palette="bright")
plt.show()
```



## 0.1 Sklearn routine

- Initialize method
- fit
- transform/predict

### 0.1.1 Normalizations

```
[177]: # z-score, fit and then transform
```

```
#Standardize features by removing the mean and scaling to unit variance.
```

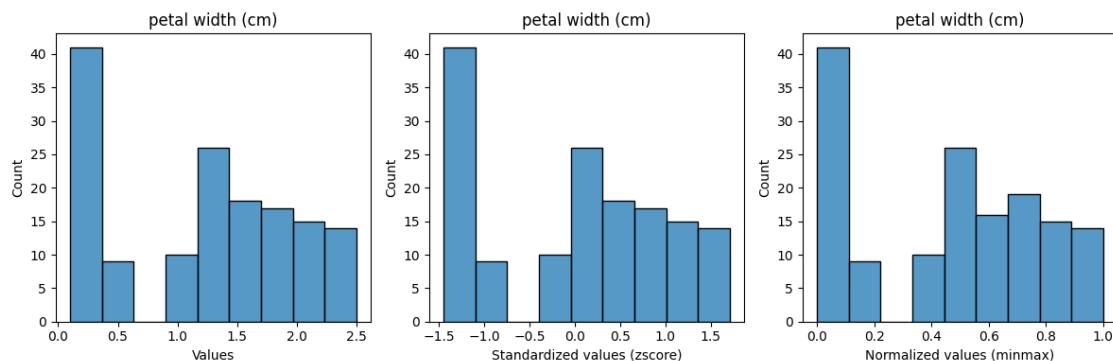
```
# z = (x - u) / s

# u is the mean and s the standard deviation

scaler = StandardScaler()
scaler.fit(X)
X_scal = scaler.transform(X)
```

[178]: # min-max, fit and transform directly  
scaler = MinMaxScaler()  
X\_minmax = scaler.fit\_transform(X)

[179]: i = 3 # column index  
fig, axs = plt.subplots(1,3, figsize=(12, 4)) # 1 row, 3 columns  
  
sns.histplot(X[:,i], ax=axs[0]).set(title=df.columns[i])
axs[0].set(xlabel='Values')  
  
sns.histplot(X\_scal[:,i], ax=axs[1]).set(title=df.columns[i])
axs[1].set(xlabel='Standardized values (zscore)')  
  
sns.histplot(X\_minmax[:,i], ax=axs[2]).set(title=df.columns[i])
axs[2].set(xlabel='Normalized values (minmax)')  
  
plt.tight\_layout() # Adjust the padding between and around subplots



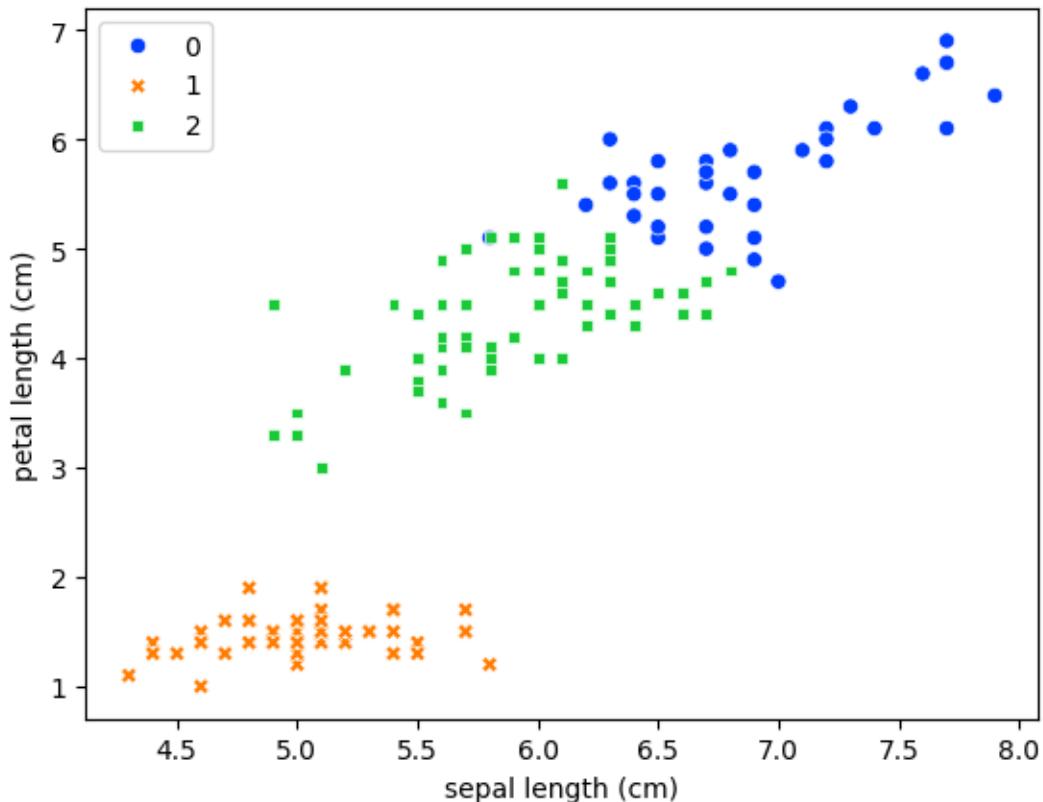
### 0.1.2 Kmeans

[180]: kmeans = KMeans(n\_clusters=3, n\_init=10, max\_iter=100, random\_state=94)  
kmeans.fit(X\_minmax)

[180]: KMeans(max\_iter=100, n\_clusters=3, n\_init=10, random\_state=94)

[181]: kmeans.labels\_

```
[182]: sns.scatterplot(data=df,
                      x="sepal length (cm)",
                      y="petal length (cm)",
                      hue=kmeans.labels_,
                      style=kmeans.labels_,
                      palette="bright")
plt.show()
```



```
[ ]: # ndarray of shape (n_clusters, n_features)
      kmeans.cluster_centers_
```

```
[ ]: array([[0.70726496, 0.4508547 , 0.79704476, 0.82478632],  
          [0.19611111, 0.595      , 0.07830508, 0.06083333],
```

```
[0.44125683, 0.30737705, 0.57571548, 0.54918033]])
```

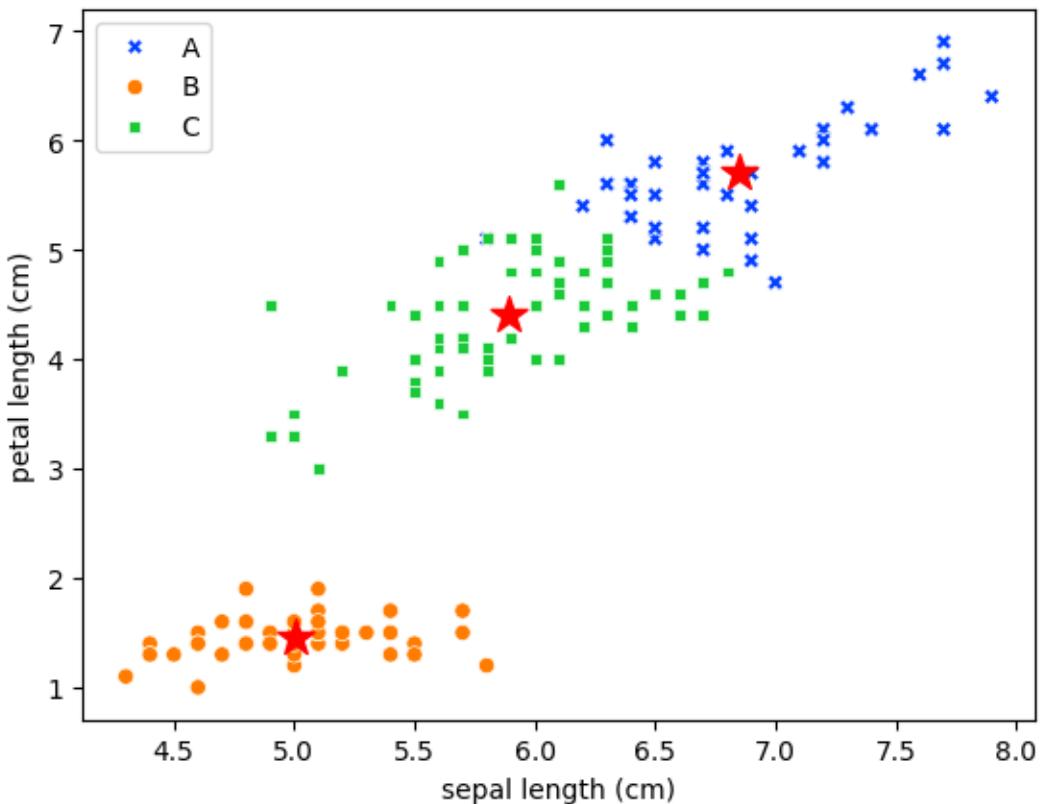
```
[ ]: centers = scaler.inverse_transform(kmeans.cluster_centers_)

centers
```

```
[ ]: array([[6.84615385, 3.08205128, 5.7025641 , 2.07948718],
           [5.006      , 3.428      , 1.462      , 0.246      ],
           [5.88852459, 2.73770492, 4.39672131, 1.41803279]])
```

```
[ ]: df['kmeans_labels'] = kmeans.labels_
df['kmeans_labels'] = df['kmeans_labels'].map({0:"A", 1: "B", 2: "C"})
```

```
[ ]: sns.scatterplot(data=df,
                     x="sepal length (cm)",
                     y="petal length (cm)",
                     hue='kmeans_labels',
                     style='kmeans_labels',
                     palette="bright",
                     hue_order=["A", "B", "C"]
)
plt.legend()
plt.scatter(centers[:,0], centers[:,2], c='red', marker='*', s=200)
plt.show()
```



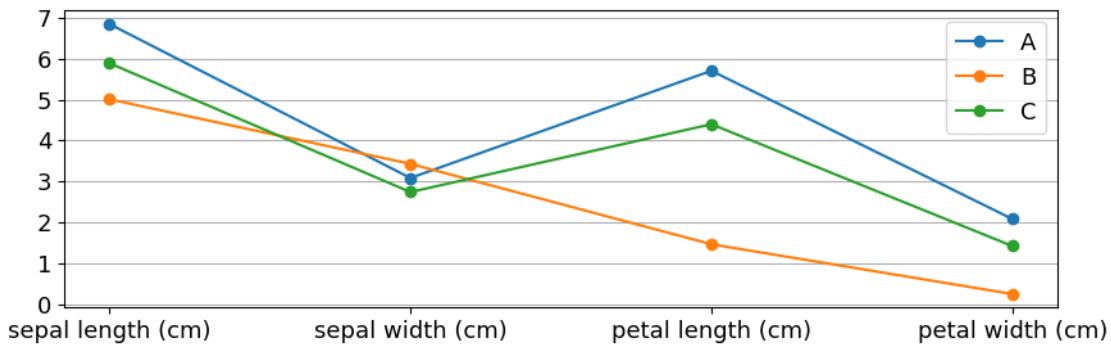
```
[ ]: plt.figure(figsize=(10, 3))

clust_name = ['A', 'B', 'C']

for i in range(len(centers)):
    plt.plot(centers[i], marker='o', label=clust_name[i])

plt.xticks(range(0, len(df.columns) - 1), df.columns[:-1], fontsize=13)
plt.yticks(fontsize=13)

plt.legend(fontsize=13, loc='best')
plt.grid(axis='y')
```



```
[ ]: plt.figure(figsize=(10, 3))

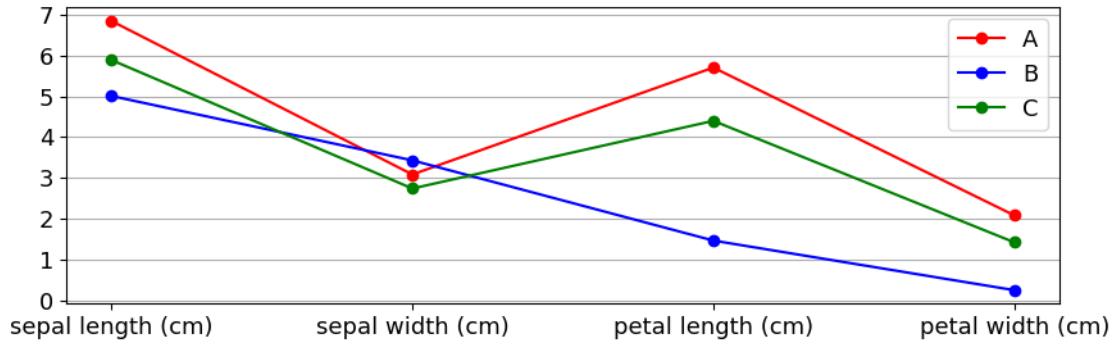
clust_name = ['A', 'B', 'C']

# color customization (fixed to 3 colors, it would not work if the clusters ↴ would be 5)
colors = {
    'A': 'r',
    'B': 'b',
    'C': 'g',
}

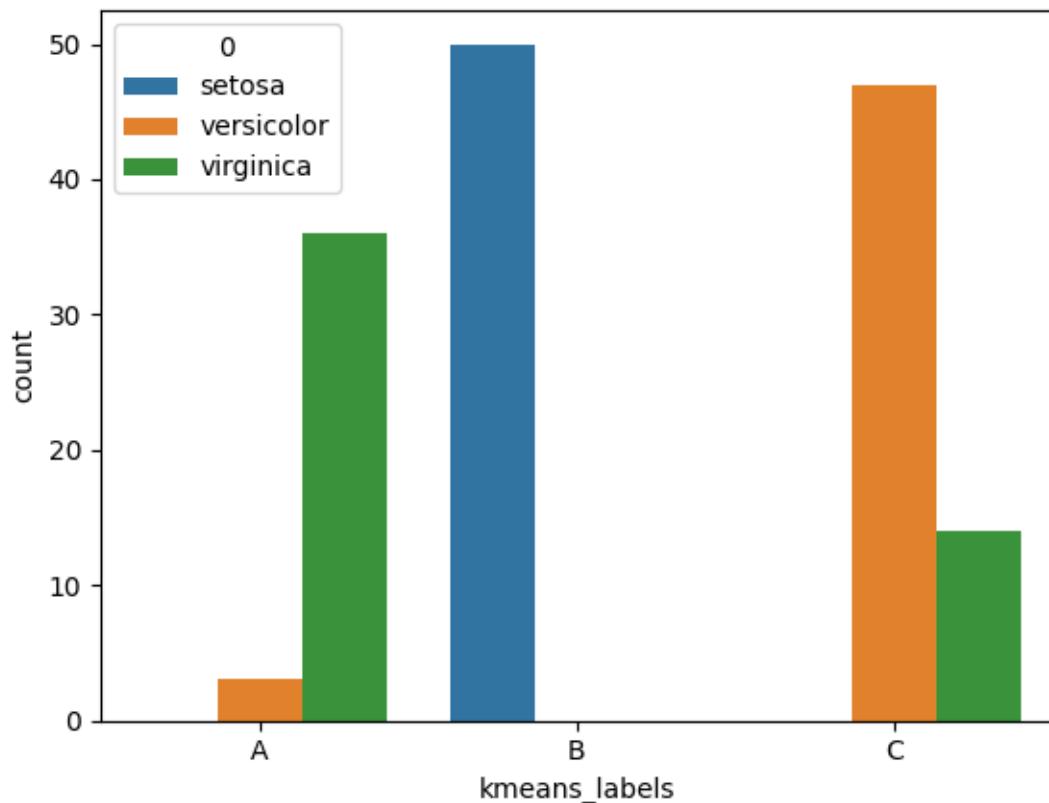
for i in range(len(centers)):
    plt.plot(centers[i], marker='o', label=clust_name[i], color=colors[clust_name[i]])

plt.xticks(range(0, len(df.columns) - 1), df.columns[:-1], fontsize=13)
plt.yticks(fontsize=13)
```

```
plt.legend(fontsize=13, loc='best')
plt.grid(axis='y')
```



```
[ ]: sns.countplot(data=df, x='kmeans_labels', hue=y_mapped, order=["A", "B", "C"])
plt.show()
```



```
[ ]: y_mapped
```

```
[ ]: 0      setosa
1      setosa
2      setosa
3      setosa
4      setosa
...
145    virginica
146    virginica
147    virginica
148    virginica
149    virginica
Name: 0, Length: 150, dtype: object
```

```
[ ]: y_mapped.groupby(df['kmeans_labels']).value_counts(normalize=True)
```

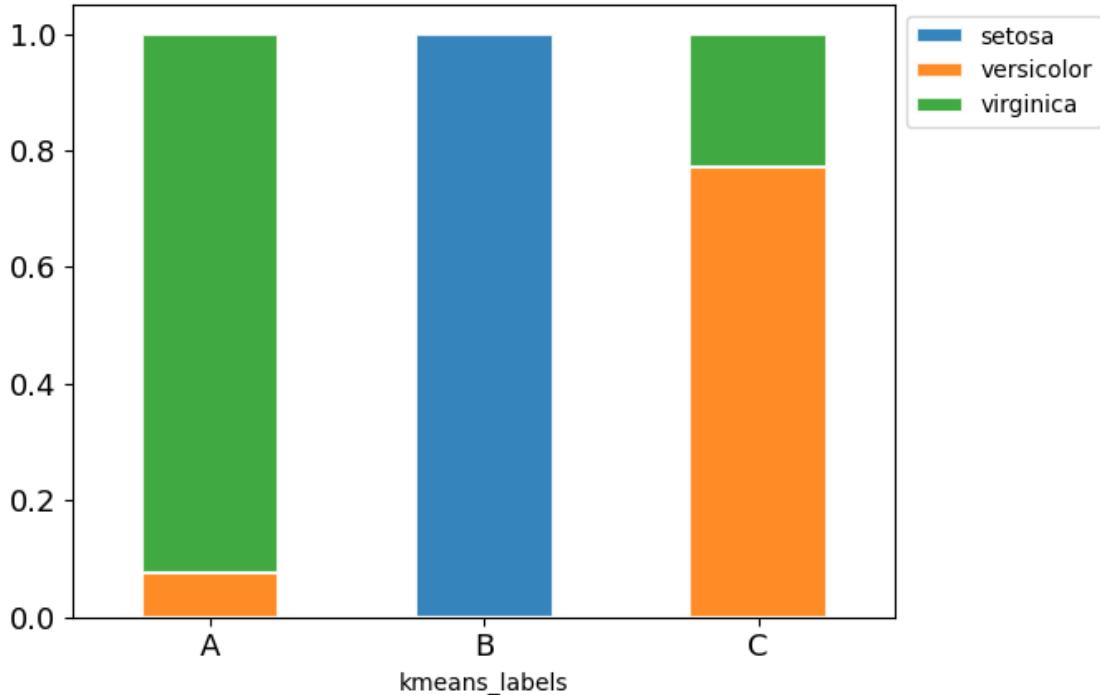
```
[ ]: kmeans_labels  0
A            virginica    0.923077
              versicolor   0.076923
B            setosa     1.000000
C            versicolor   0.770492
              virginica    0.229508
Name: proportion, dtype: float64
```

```
[ ]:
```

```
[ ]: bar_pl = y_mapped.groupby(df['kmeans_labels']).value_counts(normalize=True).
     ↪unstack(1)
bar_pl.plot(kind='bar', stacked=True, alpha=0.9, edgecolor='white', linewidth=1.
     ↪5)

plt.xticks(range(0, len(clust_name)), clust_name, fontsize=13, rotation=0)
plt.yticks(fontsize=13)
plt.legend(bbox_to_anchor=(1,1))

plt.show()
```



```
[ ]: np.unique(kmeans.labels_, return_counts=True)
```

```
[ ]: (array([0, 1, 2], dtype=int32), array([39, 50, 61]))
```

```
[ ]: print('SSE', kmeans.inertia_)
print('Silhouette', silhouette_score(X_minmax, kmeans.labels_))
```

SSE 6.982216473785234  
 Silhouette 0.5047687565398589

```
[ ]: # how do we select the number of clusters k?
```

```
[ ]: %%time
sse_list = []

for k in range(2, 51):
    kmeans = KMeans(n_clusters=k, n_init=10, max_iter=100, random_state = 42)
    kmeans.fit(X_minmax)
    sse_list.append(kmeans.inertia_)
```

CPU times: user 4.39 s, sys: 19 ms, total: 4.41 s  
 Wall time: 6.84 s

```
[ ]: sse_list
```

[ ]: [12.12779075053819,  
6.982216473785234,  
5.516933472040375,  
4.5809486401172945,  
3.97642424820984,  
3.473622320733247,  
3.1456415905967785,  
2.814016644149028,  
2.5608613580326587,  
2.2987533401137896,  
2.164318065645078,  
2.0022593438158807,  
1.9136218419482227,  
1.845344088641389,  
1.7250142667163755,  
1.6428239677141323,  
1.5736603125876003,  
1.4746274728453737,  
1.3854324519788492,  
1.2910389408577807,  
1.2772755130523463,  
1.1940938985455984,  
1.1577934403006787,  
1.1071461465584955,  
1.0704311954625618,  
1.0267934761861963,  
0.9752628007631279,  
0.9344469529963271,  
0.9098453553268242,  
0.898175535380677,  
0.8211419791349482,  
0.808951770629316,  
0.7788104386486097,  
0.7628687044231548,  
0.7704024908892751,  
0.715901256957508,  
0.6921080400722284,  
0.6728572465814967,  
0.632164669967409,  
0.6184770195918264,  
0.6067667847598058,  
0.5760920060753081,  
0.5591060283729715,  
0.518060741783396,  
0.5403486678437852,  
0.5265759718954294,  
0.48127059027582497,

```
0.470708307743383,  
0.4707451756404255]
```

```
[ ]: %%time  
sil_list = []  
  
for k in range(2, 51):  
    kmeans = KMeans(n_clusters=k, n_init=10, max_iter=100, random_state = 42)  
    kmeans.fit(X_minmax)  
    sil_list.append(silhouette_score(X_minmax, kmeans.labels_))
```

```
CPU times: user 3.71 s, sys: 16.6 ms, total: 3.72 s  
Wall time: 3.81 s
```

```
[ ]: pdist(X_minmax).shape
```

```
[ ]: (11175,)
```

```
[ ]: squareform(pdist(X_minmax))
```

```
[ ]: array([[0.          , 0.21561354, 0.16810102, ... , 1.08257132, 1.14907064,  
           0.96462829],  
          [0.21561354, 0.          , 0.10157824, ... , 1.08390691, 1.17619813,  
           0.95649502],  
          [0.16810102, 0.10157824, 0.          , ... , 1.12088708, 1.19544459,  
           0.98859665],  
          ... ,  
          [1.08257132, 1.08390691, 1.12088708, ... , 0.          , 0.226928 ,  
           0.18710825],  
          [1.14907064, 1.17619813, 1.19544459, ... , 0.226928 , 0.          ,  
           0.28409587],  
          [0.96462829, 0.95649502, 0.98859665, ... , 0.18710825, 0.28409587,  
           0.        ]])
```

```
[ ]: %%time  
sil_list = []  
dist = squareform(pdist(X_minmax)) # using a precomputed distance matrix  
  
for k in range(2, 51):  
    kmeans = KMeans(n_clusters=k, n_init=10, max_iter=100)  
    kmeans.fit(X_minmax)  
    sil_list.append(silhouette_score(dist, kmeans.labels_,  
                                     metric='precomputed'))
```

```
CPU times: user 2.85 s, sys: 15.3 ms, total: 2.86 s  
Wall time: 2.88 s
```

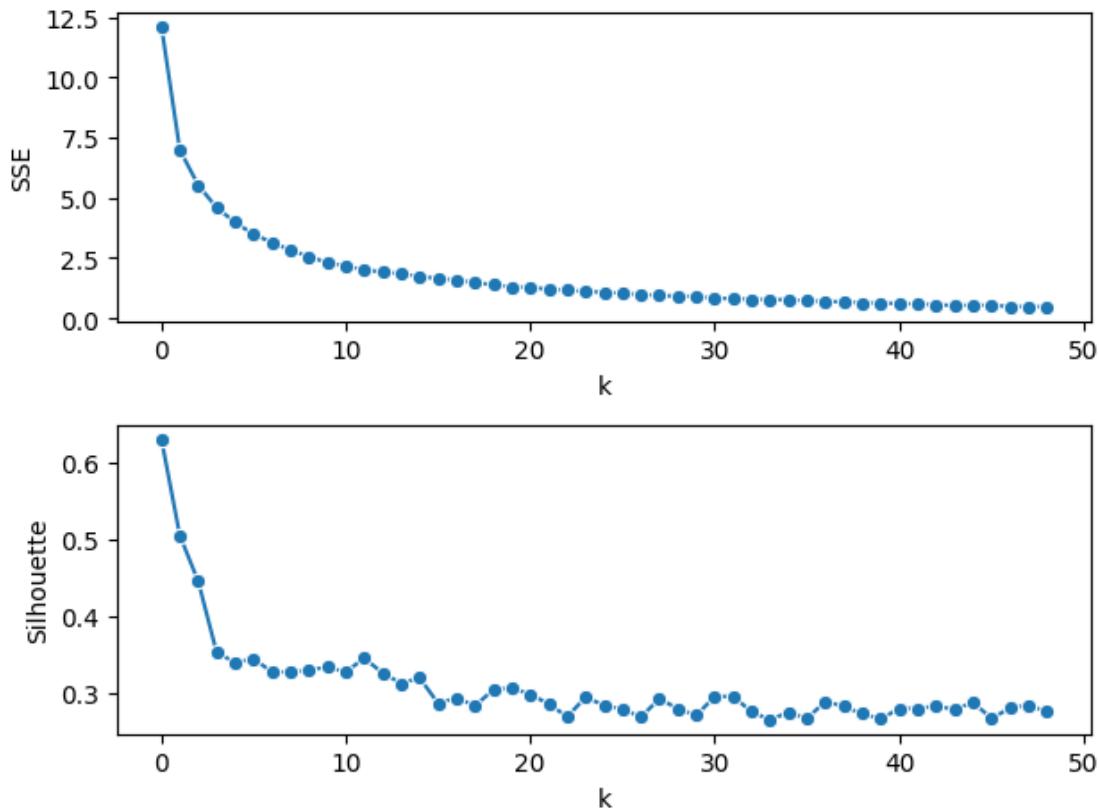
```
[ ]:
```

```
[ ]: fig, axs = plt.subplots(2)

sns.lineplot(x=range(len(sse_list)), y=sse_list, marker='o', ax=axs[0])
axs[0].set(xlabel='k', ylabel='SSE')

sns.lineplot(x=range(len(sil_list)), y=sil_list, marker='o', ax=axs[1])
axs[1].set(xlabel='k', ylabel='Silhouette')

plt.tight_layout() # Adjust the padding between and around subplots
```



### 0.1.3 Bisecting K-means

```
[ ]: from sklearn.cluster import BisectingKMeans

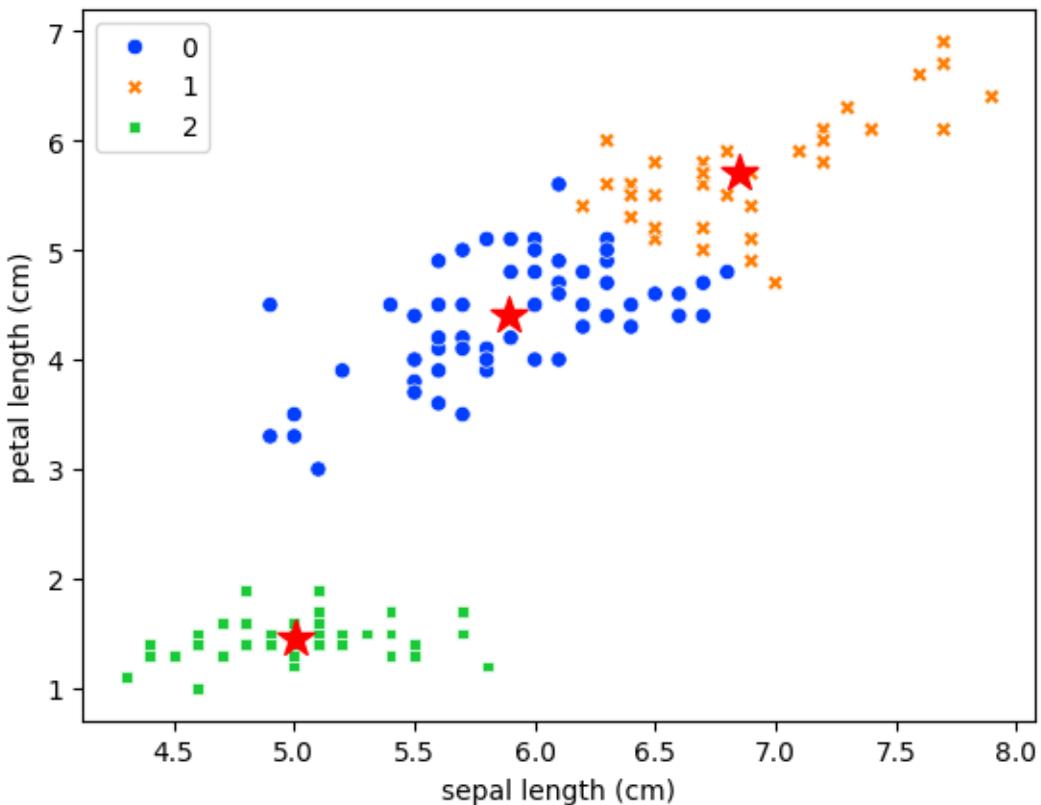
[ ]: bkmeans = BisectingKMeans(n_clusters=3, n_init=10, max_iter=100, random_state=94)
      bkmeans.fit(X_minmax)

[ ]: BisectingKMeans(max_iter=100, n_clusters=3, n_init=10, random_state=94)
```

```
[ ]: centers = scaler.inverse_transform(bkmeans.cluster_centers_)
centers
```

```
[ ]: array([[5.88852459, 2.73770492, 4.39672131, 1.41803279],
       [6.84615385, 3.08205128, 5.7025641 , 2.07948718],
       [5.006      , 3.428      , 1.462      , 0.246      ]])
```

```
[ ]: sns.scatterplot(data=df,
                     x="sepal length (cm)",
                     y="petal length (cm)",
                     hue=bkmeans.labels_,
                     style=bkmeans.labels_,
                     palette="bright",
                     )
plt.legend()
plt.scatter(centers[:,0], centers[:,2], c='red', marker='*', s=200)
plt.show()
```



```
[ ]: print('SSE', bkmeans.inertia_)
print('Silhouette', silhouette_score(X_minmax, bkmeans.labels_))
```

SSE 6.982216473785234

```
Silhouette 0.5047687565398589
```

```
[ ]: # distance matrix to speed up sil score
```

#### 0.1.4 k medoids

```
[ ]: #!pip install pyclustering
```

```
[ ]: # standard installation might result in error due to numpy warnings (numpy > 1.  
→24.0)
```

```
#after installing pyclustering, also install this warning fix below
```

```
#!pip install https://github.com/KulikDM/pyclustering/archive/Warning-Fix.zip
```

```
[183]: from pyclustering.cluster import kmedoids
```

```
[184]: # Set random initial medoids.
```

```
initial_medoids = [1, 50]
```

```
kmedoids_instance = kmedoids.kmedoids(X, initial_medoids)  
kmedoids_instance.process()
```

```
[184]: <pyclustering.cluster.kmedoids.kmedoids at 0x7dd7f437f9e0>
```

```
[185]: clusters = kmedoids_instance.get_clusters()
```

```
[186]: clusters = [np.array(x) for x in clusters]
```

```
[187]: clusters3
```

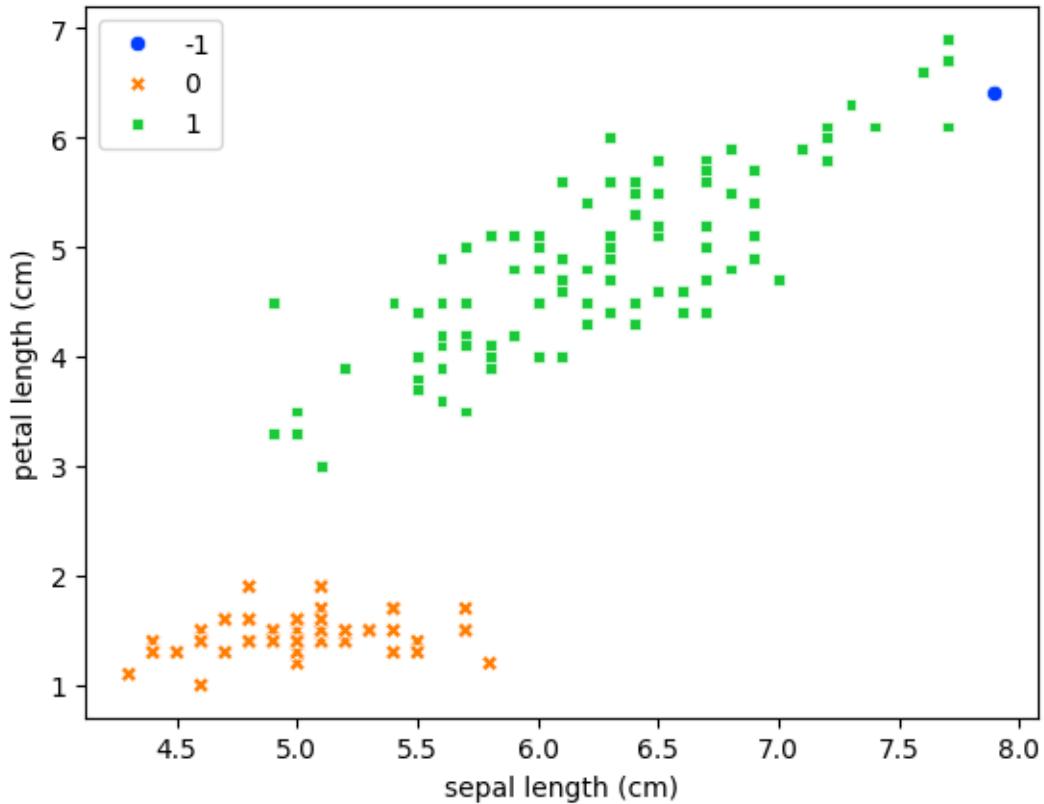
```
[187]: [array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,  
       17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,  
       34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 57,  
       98]),  
 array([ 50,  51,  52,  53,  54,  55,  56,  58,  59,  60,  61,  62,  63,  
       64,  65,  66,  67,  68,  69,  70,  71,  72,  73,  74,  75,  76,  
       77,  78,  79,  80,  81,  82,  83,  84,  85,  86,  87,  88,  89,  
       90,  91,  92,  93,  94,  95,  96,  97,  99, 100, 101, 102, 103,  
      104, 105, 106, 107, 108, 109, 110, 111, 112, 113, 114, 115, 116,  
      117, 118, 119, 120, 121, 122, 123, 124, 125, 126, 127, 128, 129,  
      130, 131, 132, 133, 134, 135, 136, 137, 138, 139, 140, 141, 142,  
      143, 144, 145, 146, 147, 148, 149])]
```

### 0.1.5 DBScan

```
[ ]: dbSCAN = DBSCAN(eps=0.3, min_samples=5)
dbSCAN.fit(X_minmax)
```

```
[ ]: DBSCAN(eps=0.3)
```

```
[ ]: sns.scatterplot(data=df,
                     x="sepal length (cm)",
                     y="petal length (cm)",
                     hue=dbSCAN.labels_,
                     style=dbSCAN.labels_,
                     palette="bright")
plt.show()
```



```
[ ]: np.unique(dbSCAN.labels_, return_counts=True)
```

```
[ ]: (array([-1,  0,  1]), array([ 1, 50, 99]))
```

```
[ ]: dbSCAN.labels_
```

```
[ ]: print('Silhouette', silhouette_score(X_minmax, dbscan.labels_)) # counting  
    ↳silhouette also w.r.t to noise cluster (-1)  
print('Silhouette', silhouette_score(X_minmax[dbscan.labels_ != -1], dbscan.  
    ↳labels_[dbscan.labels_ != -1]))
```

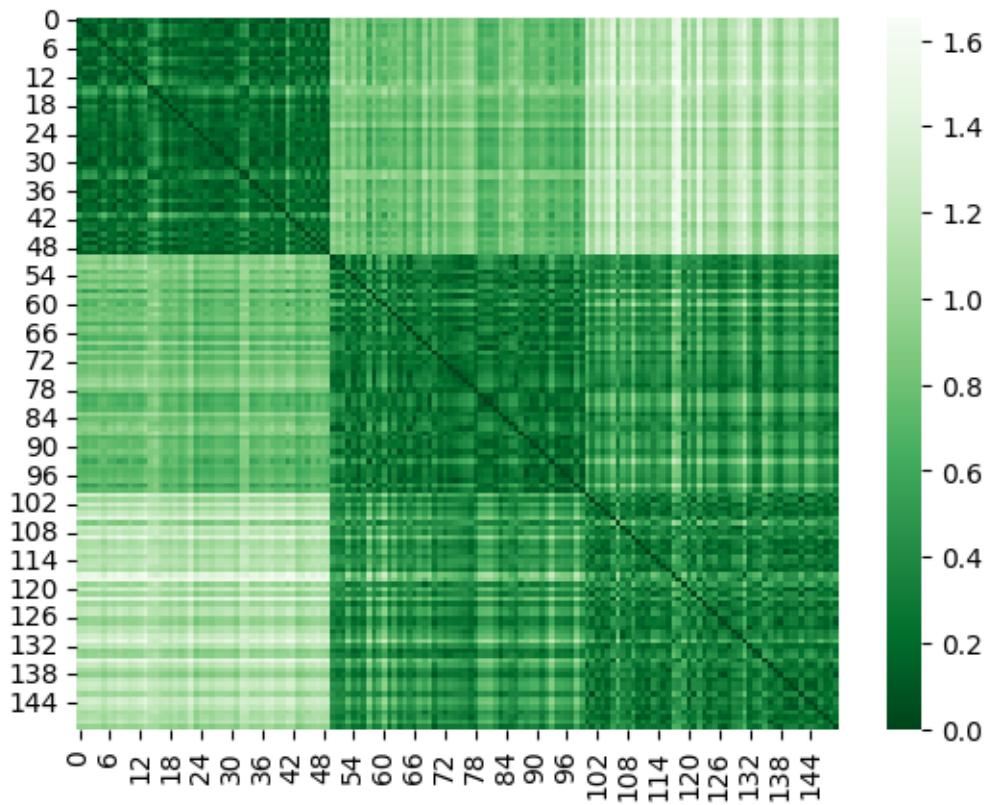
Silhouette 0.4681853590627469

Silhouette 0.6331315097272541

### 0.1.6 kth neighbor distance

```
[ ]: dist = pdist(X_minmax, 'euclidean')
      dist = squareform(dist)
```

```
[ ]: sns.heatmap(dist, cmap="Greens_r", annot=False)
      plt.show()
```

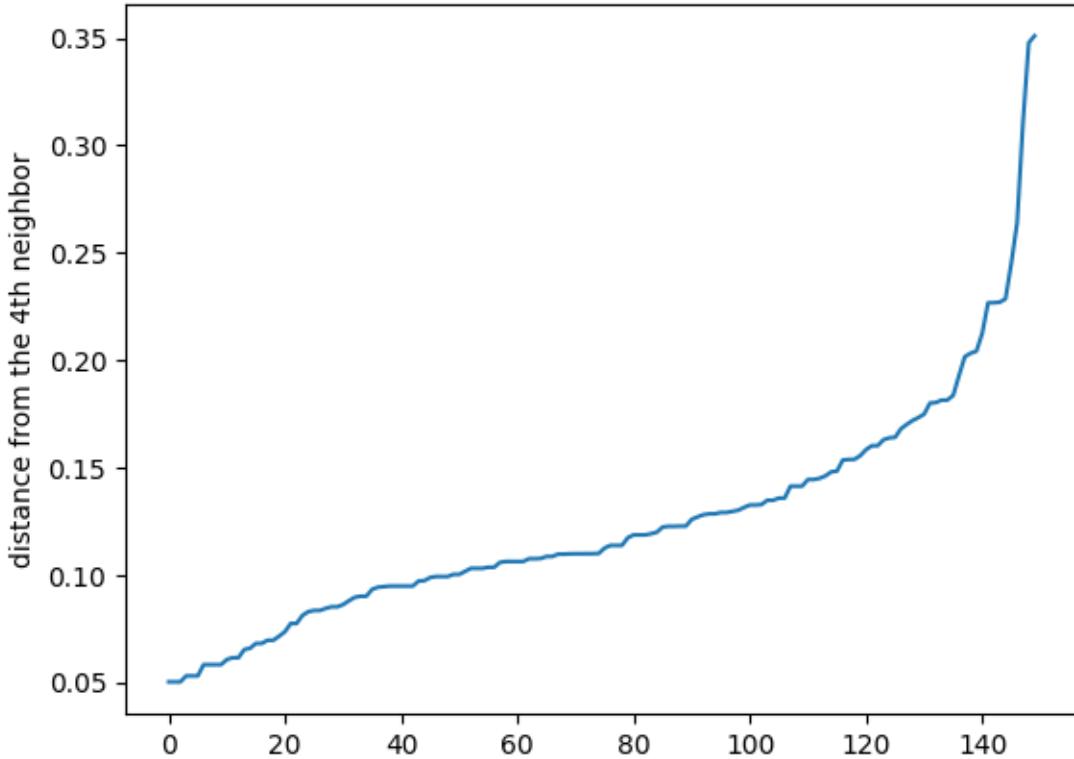


```
[ ]: k=4
kth_distances = []

for d in dist: # d is a vector containing distances between the ith record and
    # all the others
    index_kth_distance = np.argsort(d)[k] # take the index of the kth nearest
    # neighbor
    kth_distances.append(d[index_kth_distance]) # store the distance in a list

[ ]: plt.plot(range(0, len(kth_distances)), sorted(kth_distances))
plt.ylabel('distance from the {}th neighbor'.format(k))

plt.show()
```



```
[ ]: ks = [4, 8, 16, 32, 62]

fig = plt.figure(figsize=(16, 3)) # dimensions of the overall plot
fig_dims = (1, len(ks))

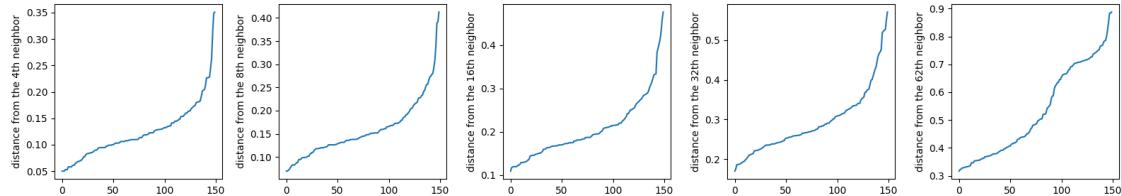
for i in range(len(ks)):
    k = ks[i]
    kth_distances = []

    for d in dist: # d is a vector containing distances between the ith record
        ↪and all the others
        index_kth_distance = np.argsort(d)[k] # take the index of the kth
        ↪nearest neighbor
        kth_distances.append(d[index_kth_distance]) # store the distance in a
        ↪list

    plt.subplot2grid(fig_dims, (0, i))
    plt.plot(range(0, len(kth_distances)), sorted(kth_distances))
    plt.ylabel('distance from the {}th neighbor'.format(k))

plt.tight_layout()
```

```
plt.show()
```

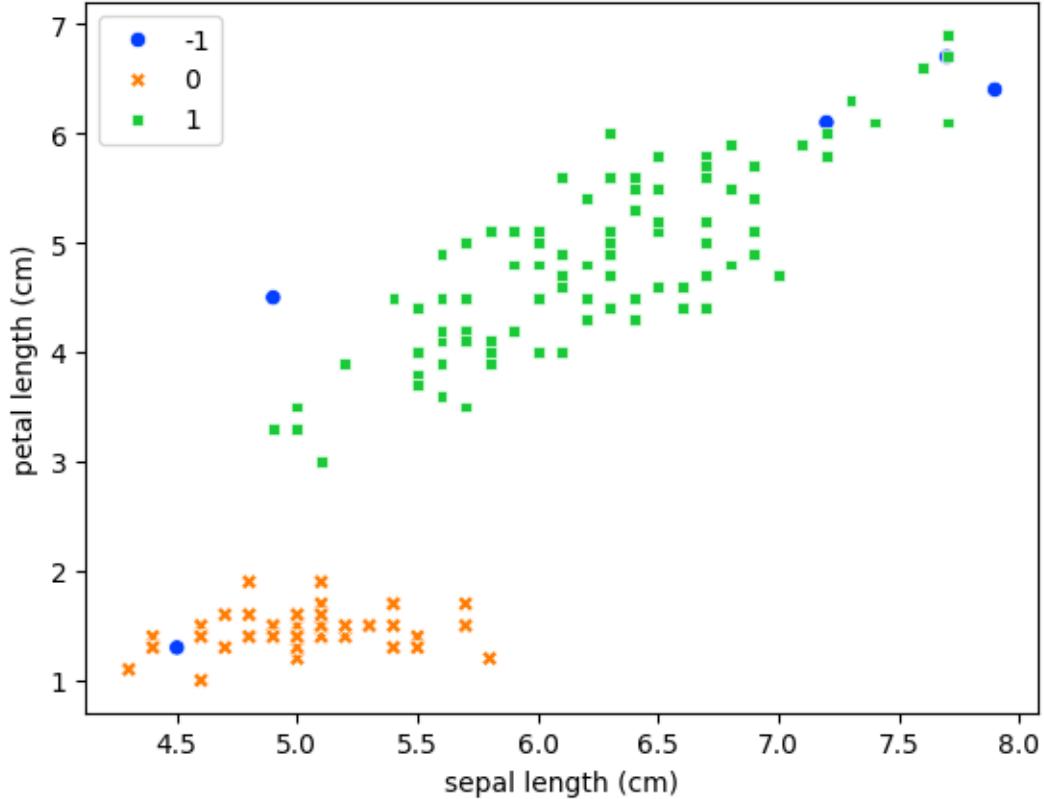


```
[ ]: dbSCAN = DBSCAN(eps=0.19, min_samples=4, metric='precomputed')
dbSCAN.fit(dist)
```

```
[ ]: DBSCAN(eps=0.19, metric='precomputed', min_samples=4)
```

```
[ ]: sns.scatterplot(data=df,
                     x="sepal length (cm)",
                     y="petal length (cm)",
                     hue=dbSCAN.labels_,
                     style=dbSCAN.labels_,
                     palette="bright")
```

```
plt.show()
```



```
[ ]: np.unique(dbSCAN.labels_, return_counts=True)
[ ]: (array([-1,  0,  1]), array([ 5, 49, 96]))
[ ]: df['dblables'] = dbSCAN.labels_
df.head()

[ ]:      sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm) \
0                 5.1           3.5             1.4            0.2
1                 4.9           3.0             1.4            0.2
2                 4.7           3.2             1.3            0.2
3                 4.6           3.1             1.5            0.2
4                 5.0           3.6             1.4            0.2

      kmeans_labels  dblables
0               B          0
1               B          0
2               B          0
3               B          0
4               B          0

[ ]: # dbSCAN centroids
df[['sepal length (cm)', 'sepal width (cm)', 'petal width (cm)', 'petal length (cm)', 'dblables']].groupby('dblables').mean()

[ ]:      sepal length (cm)  sepal width (cm)  petal width (cm) \
dblables
-1                 6.440000           3.200000           1.740000
0                  5.016327           3.451020           0.244898
1                  6.234375           2.848958           1.658333

      petal length (cm)
dblables
-1                 5.000000
0                  1.465306
1                  4.863542

[ ]: print('Silhouette', silhouette_score(X_minmax[dbSCAN.labels_ != -1], dbSCAN.
                                         labels_[dbSCAN.labels_ != -1]))
Silhouette 0.6432450335271236
### OPTICS
[ ]: from sklearn.cluster import OPTICS, cluster_optics_dbSCAN
```

```
[ ]: optics = OPTICS(min_samples = 4, max_eps = np.inf, min_cluster_size=40)
optics.fit(X_minmax)

[ ]: OPTICS(min_cluster_size=40, min_samples=4)

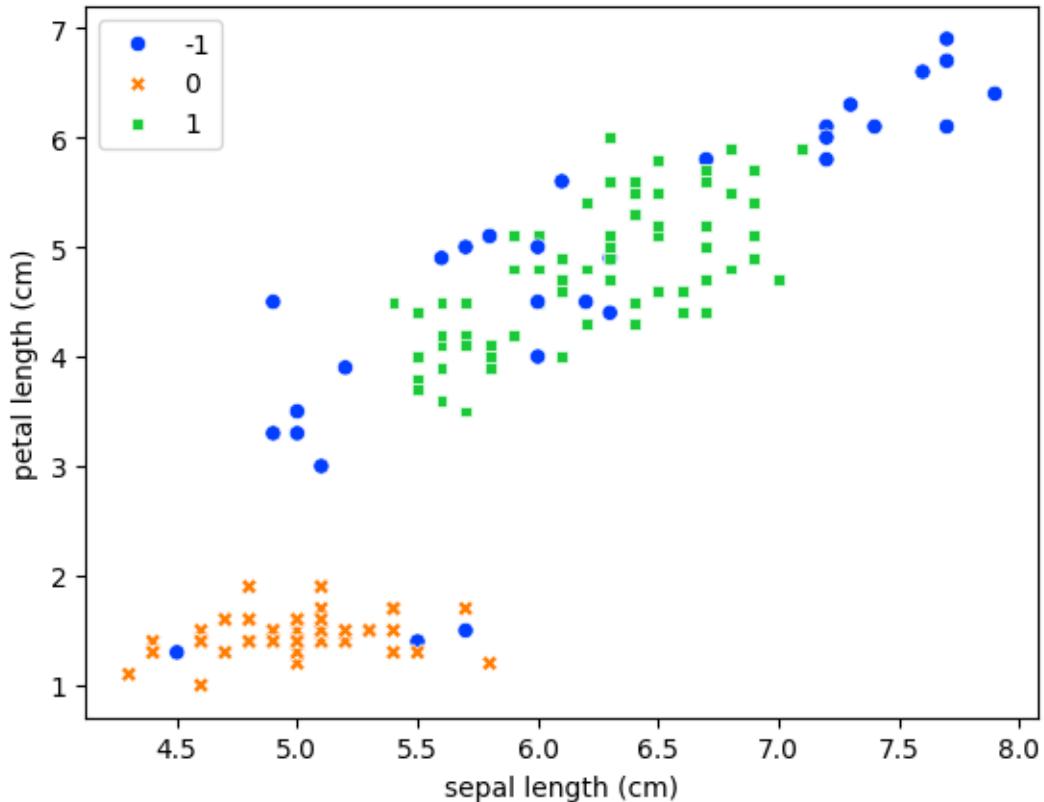
[ ]: print("Silhouette", silhouette_score(X_minmax[optics.labels_ != -1], optics.
                                          labels_[optics.labels_ != -1]))

Silhouette 0.6901722966689927

[ ]: np.unique(optics.labels_, return_counts=True)

[ ]: (array([-1,  0,  1]), array([33, 47, 70]))

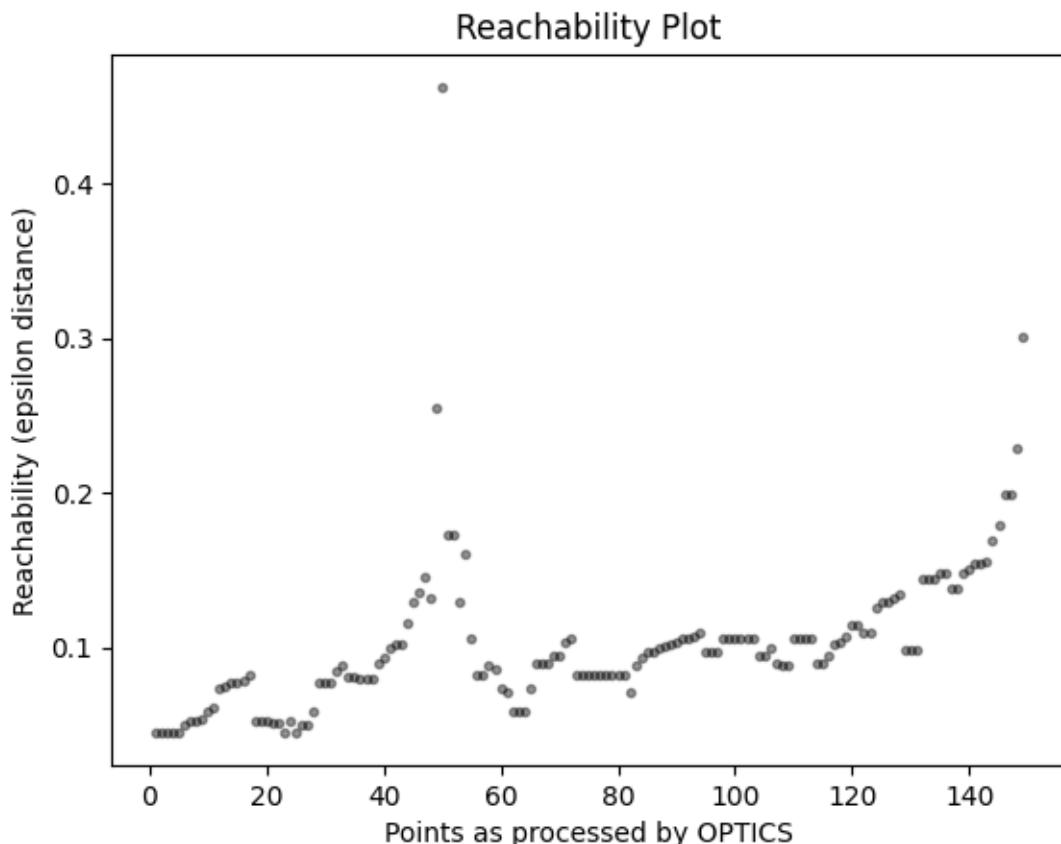
[ ]: sns.scatterplot(data=df,
                     x="sepal length (cm)",
                     y="petal length (cm)",
                     hue=optics.labels_,
                     style=optics.labels_,
                     palette="bright")
plt.show()
```



```
[ ]: # Valleys in the plot correspond to the clusters
space = np.arange(len(X_minmax))
reachability = optics.reachability_[optics.ordering_]
labels = optics.labels_[optics.ordering_]

# Reachability plot
plt.plot(space, reachability, "k.", alpha=0.4)
plt.xlabel("Points as processed by OPTICS")
plt.ylabel("Reachability (epsilon distance)")
# plt.axhline(0.3, c="red")
plt.title("Reachability Plot")
```

```
[ ]: Text(0.5, 1.0, 'Reachability Plot')
```



### 0.1.7 Hierarchical

```
[ ]: def get_linkage_matrix(model):
    # Create linkage matrix

    # create the counts of samples under each node
```

```

counts = np.zeros(model.children_.shape[0])
n_samples = len(model.labels_)
for i, merge in enumerate(model.children_):
    current_count = 0
    for child_idx in merge:
        if child_idx < n_samples:
            current_count += 1 # leaf node
        else:
            current_count += counts[child_idx - n_samples]
    counts[i] = current_count

linkage_matrix = np.column_stack(
    [model.children_, model.distances_, counts]
).astype(float)

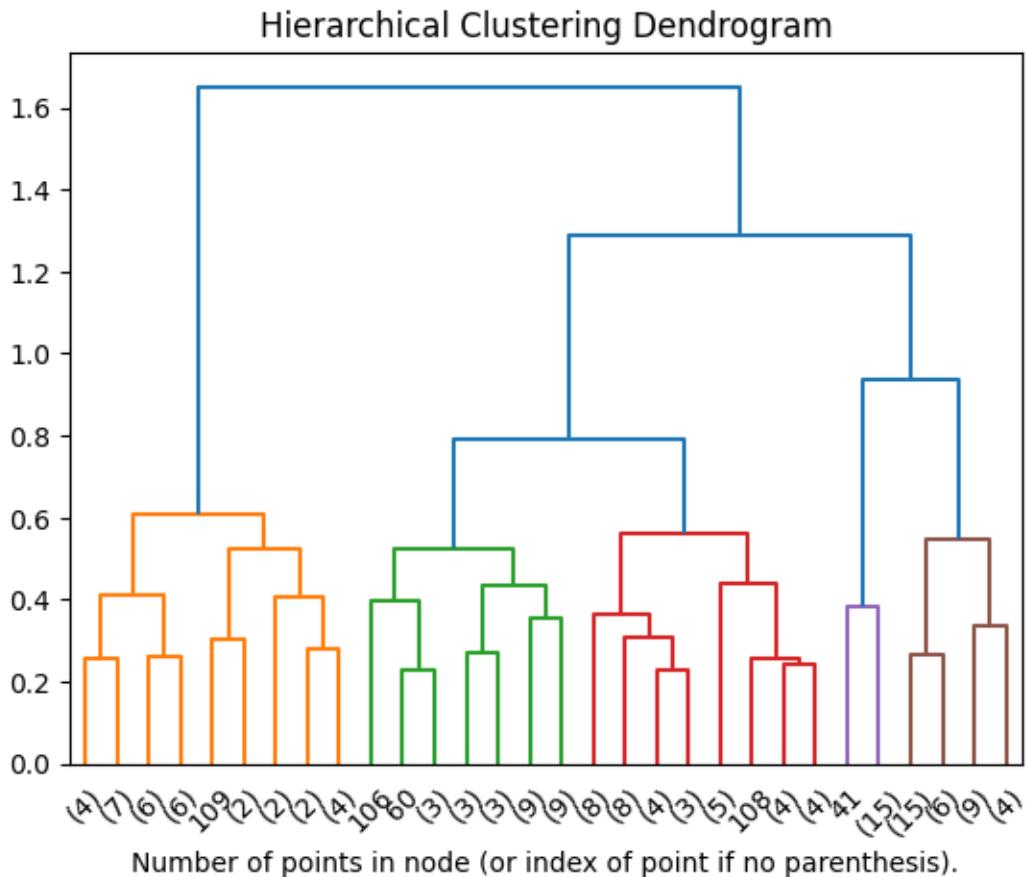
return linkage_matrix

def plot_dendrogram(model, **kwargs):
    linkage_matrix = get_linkage_matrix(model)
    dendrogram(linkage_matrix, **kwargs)

```

```
[ ]: # setting distance_threshold=0 ensures we compute the full tree.
# it is the linkage distance threshold above which clusters will not be merged
model = AgglomerativeClustering(distance_threshold=0, n_clusters=None,
                                  metric='euclidean', linkage='complete')
model = model.fit(X_minmax)
```

```
[ ]: plt.title("Hierarchical Clustering Dendrogram")
plot_dendrogram(model, truncate_mode='lastp', color_threshold=0.7)
plt.xlabel("Number of points in node (or index of point if no parenthesis).")
plt.show()
```



```
[ ]: # get the labels according to a specific threshold value cut
Z = get_linkage_matrix(model)
labels = fcluster(Z, t=0.7, criterion='distance')
```

```
[ ]: labels
```

```
[ ]: print('Silhouette', silhouette_score(X_minmax, labels))
```

Silhouette 0.3364082243047556

## Choosing the number of clusters

```
[ ]: hier = AgglomerativeClustering(n_clusters=3, metric='euclidean', linkage='complete')
hier.fit(X_minmax)

[ ]: AgglomerativeClustering(linkage='complete', n_clusters=3)

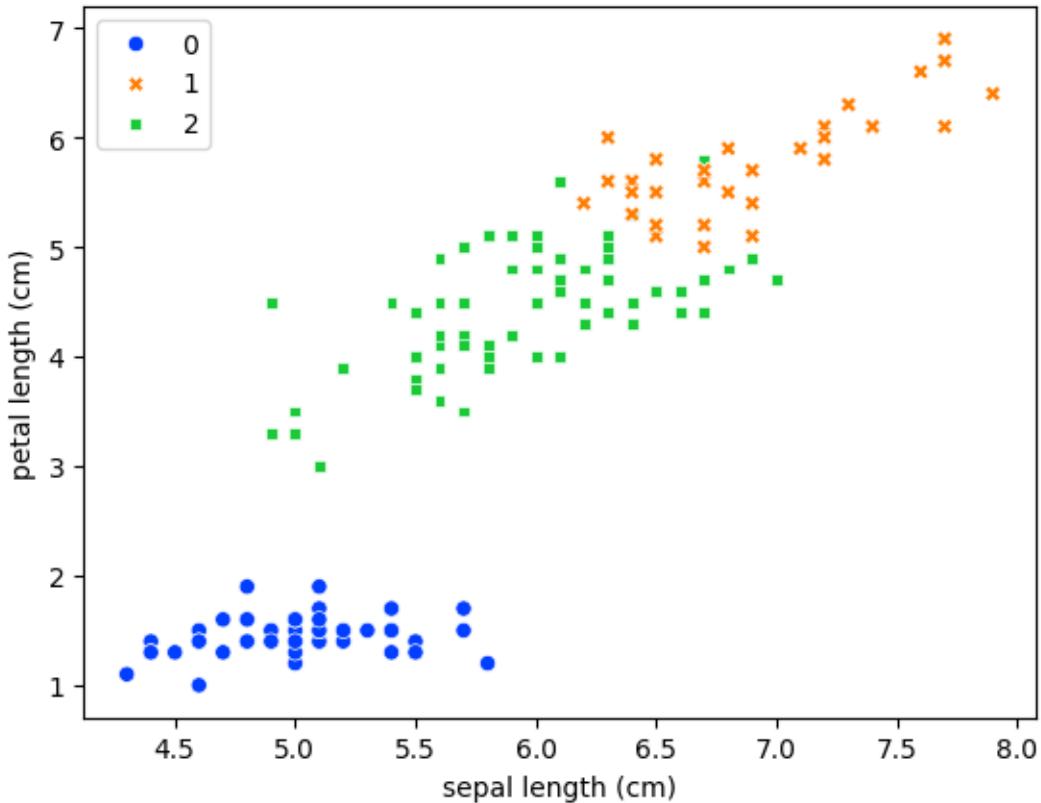
[ ]: hier.labels_

[ ]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1,
1, 2, 1, 2, 1, 1, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1,
1, 2, 2, 1, 1, 1, 2, 1, 1, 2, 1, 1, 2, 1, 1, 2, 1, 1, 2, 1, 1, 2, 1, 1, 2, 1,
1, 2, 2, 1, 1, 1, 2, 1, 1, 2, 1, 1, 2, 1, 1, 2, 1, 1, 2, 1, 1, 2, 1, 1, 2, 1])

Precomputed distance matrix
[ ]: hier = AgglomerativeClustering(n_clusters=3, metric='precomputed', linkage='complete')
hier.fit(dist)

[ ]: AgglomerativeClustering(linkage='complete', metric='precomputed', n_clusters=3)

[ ]: sns.scatterplot(data=df,
                     x="sepal length (cm)",
                     y="petal length (cm)",
                     hue=hier.labels_,
                     style=hier.labels_,
                     palette="bright")
plt.show()
```



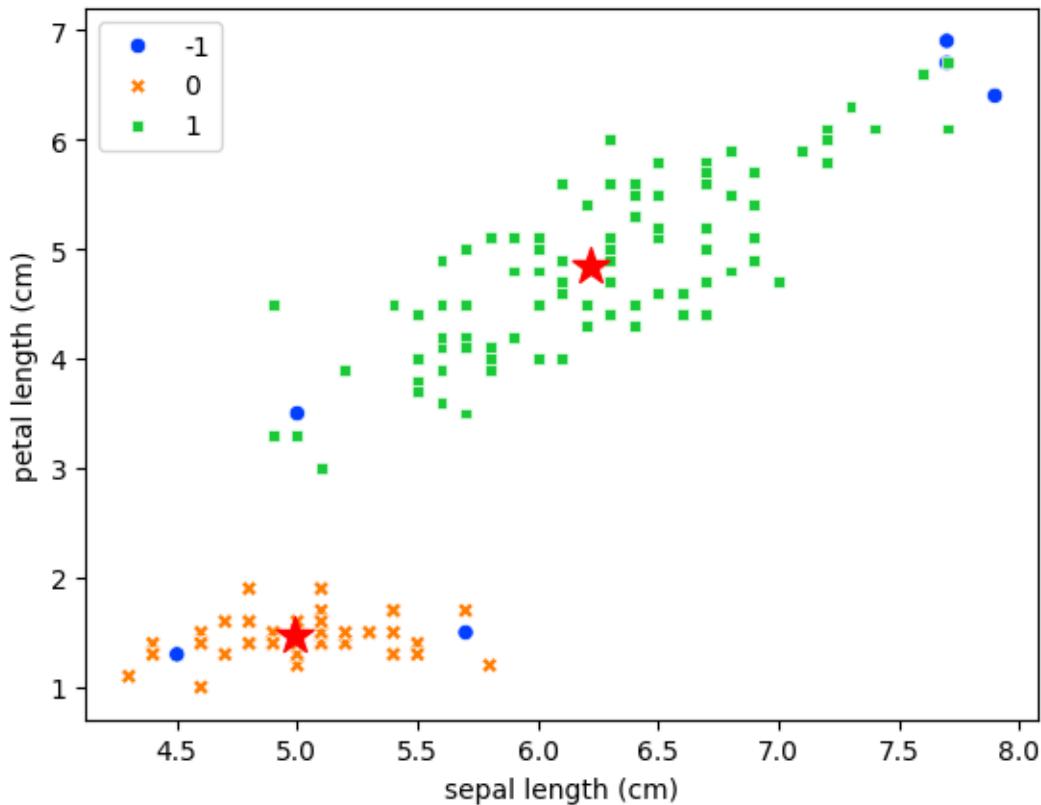
### 0.1.8 HDBScan

```
[ ]: from sklearn.cluster import HDBSCAN
[ ]: hdb = HDBSCAN(cluster_selection_epsilon=0.3, min_samples=30,
                   min_cluster_size=5, max_cluster_size=None,
                   store_centers="centroid")
hdb.fit(X_minmax)

[ ]: HDBSCAN(cluster_selection_epsilon=0.3, min_samples=30, store_centers='centroid')

[ ]: sns.scatterplot(data=df,
                     x="sepal length (cm)",
                     y="petal length (cm)",
                     hue=hdb.labels_,
                     style=hdb.labels_,
                     palette="bright")

plt.scatter(scaler.inverse_transform(hdb.centroids_)[:,0], scaler.
           inverse_transform(hdb.centroids_)[:,2], c='red', marker='*', s=200)
plt.show()
```



```
[ ]: np.unique(hdb.labels_, return_counts=True)
```

```
[ ]: (array([-1,  0,  1]), array([ 6, 48, 96]))
```

```
[ ]: hdb.centroids_
```

```
[ ]: array([[0.19207892, 0.58858965, 0.0796667 , 0.05915125],  
           [0.53191836, 0.36183505, 0.65110017, 0.64597702]])
```

```
[ ]: print("Silhouette", silhouette_score(X_minmax[hdb.labels_ != -1], hdb.  
    ↪labels [hdb.labels_ != -1]))
```

Silhouette 0.6480221541804821

[ ] :

### 9.1.9 Similarity Matrix for Cluster Validation

```
[ ]: dist = pdist(X_minmax, 'euclidean')
      dist = squareform(dist)
```

```
[ ]: hier = AgglomerativeClustering(n_clusters=3, metric='precomputed',
    ↪linkage='complete')
hier.fit(dist)

[ ]: AgglomerativeClustering(linkage='complete', metric='precomputed', n_clusters=3)

[ ]: ideal_sim

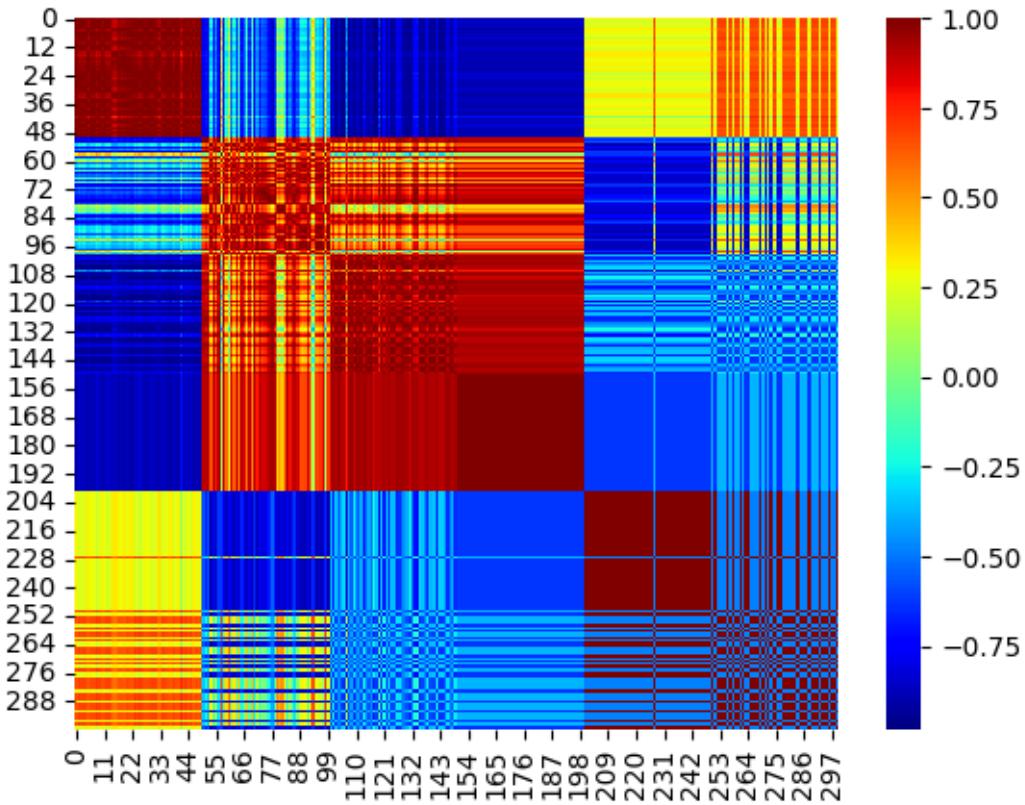
[ ]: array([[1., 1., 1., ..., 0., 0., 0.],
   [1., 1., 1., ..., 0., 0., 0.],
   [1., 1., 1., ..., 0., 0., 0.],
   ...,
   [0., 0., 0., ..., 1., 1., 0.],
   [0., 0., 0., ..., 1., 1., 0.],
   [0., 0., 0., ..., 0., 0., 1.]])
```

```
[ ]: shape = len(X_minmax)
ideal_sim = np.eye(shape)

for i in range(shape):
    for j in range(0, i+1):
        ideal_sim[j][i] = 1 if hier.labels_[i] == hier.labels_[j] else 0

ideal_sim = ideal_sim + ideal_sim.T - np.diag(np.diag(ideal_sim)) # copying
    ↪upper triangle in lower triangle
```

```
[ ]: sim = np.corrcoef(dist, ideal_sim)
sns.heatmap(sim, cmap="jet", annot=False)
plt.show()
```



### SSE Statistical evaluation

```
[ ]: import random

[ ]: kmeans = KMeans(n_clusters=3, n_init=10, max_iter=100, random_state=94)
kmeans.fit(X_minmax)
```

```
[ ]: KMeans(max_iter=100, n_clusters=3, n_init=10, random_state=94)
```

```
[ ]: my_sse = kmeans.inertia_
my_sse
```

```
[ ]: 6.982216473785234
```

```
[ ]: N = 500
sse_stats = []

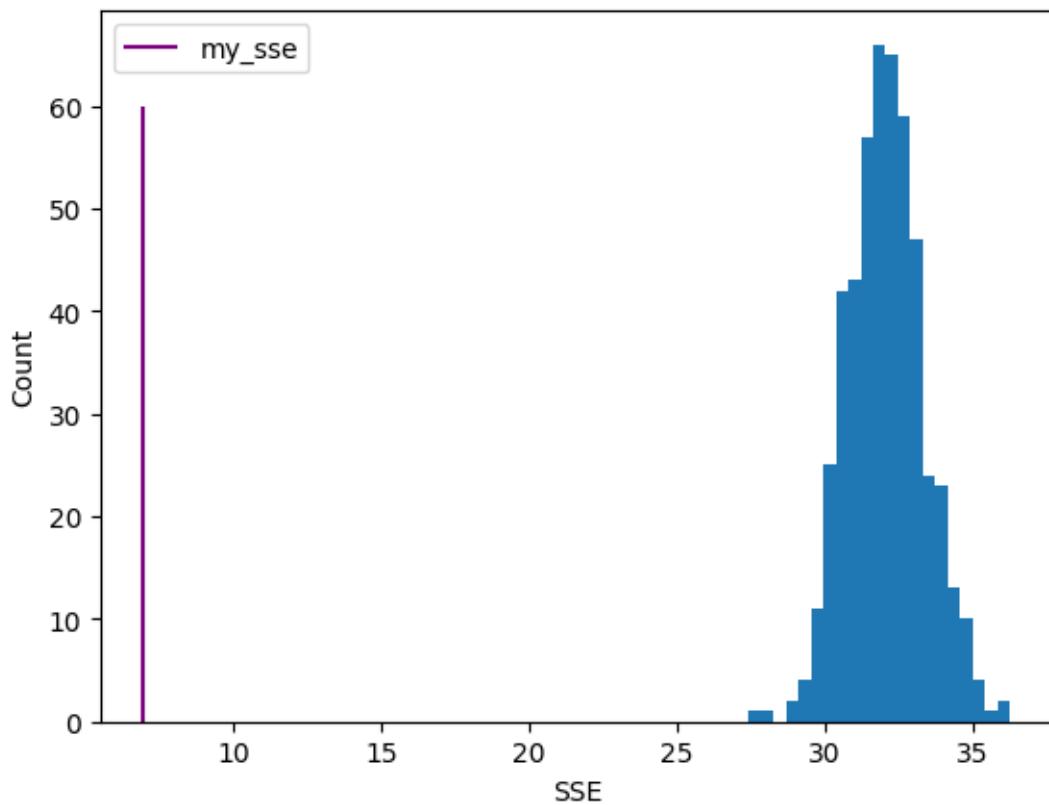
for _ in range(N):
    Xi = np.zeros(shape=X_minmax.shape)

    for cidx in range(Xi.shape[1]):
        col = X_minmax[:, cidx]
```

```
min, max, nor = np.min(col), np.max(col), len(col)
Xi[:, cidx] = np.random.uniform(min, max, (1, nor))

kmeans = KMeans(n_clusters=3, n_init=10, max_iter=100, random_state=94)
kmeans.fit(Xi)
sse_stats.append(kmeans.inertia_)
```

```
[ ]: plt.hist(sse_stats, bins='auto')
plt.vlines(x = my_sse, ymin = 0, ymax = 60, colors = 'purple', label = 'my_sse')
plt.xlabel('SSE')
plt.ylabel('Count')
plt.legend()
plt.show()
```



```
[ ]:
```