

Classification_Knn_NaiveBayes

November 22, 2025

```
[1]: !pip install scikit-plot  
      !pip install scipy==1.11.4
```

```
Requirement already satisfied: scikit-plot in /usr/local/lib/python3.12/dist-packages (0.3.7)  
Requirement already satisfied: matplotlib>=1.4.0 in /usr/local/lib/python3.12/dist-packages (from scikit-plot) (3.10.0)  
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.12/dist-packages (from scikit-plot) (1.6.1)  
Requirement already satisfied: scipy>=0.9 in /usr/local/lib/python3.12/dist-packages (from scikit-plot) (1.11.4)  
Requirement already satisfied: joblib>=0.10 in /usr/local/lib/python3.12/dist-packages (from scikit-plot) (1.5.2)  
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib>=1.4.0->scikit-plot) (1.3.3)  
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib>=1.4.0->scikit-plot) (0.12.1)  
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib>=1.4.0->scikit-plot) (4.60.1)  
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib>=1.4.0->scikit-plot) (1.4.9)  
Requirement already satisfied: numpy>=1.23 in /usr/local/lib/python3.12/dist-packages (from matplotlib>=1.4.0->scikit-plot) (1.26.4)  
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib>=1.4.0->scikit-plot) (25.0)  
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib>=1.4.0->scikit-plot) (11.3.0)  
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib>=1.4.0->scikit-plot) (3.2.5)  
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.12/dist-packages (from matplotlib>=1.4.0->scikit-plot) (2.9.0.post0)  
Requirement already satisfied: threadpoolctl>=3.1.0 in
```

/usr/local/lib/python3.12/dist-packages (from scikit-learn>=0.18->scikit-plot)
(3.6.0)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-
packages (from python-dateutil>=2.7->matplotlib>=1.4.0->scikit-plot) (1.17.0)

Requirement already satisfied: scipy==1.11.4 in /usr/local/lib/python3.12/dist-
packages (1.11.4)

Requirement already satisfied: numpy<1.28.0,>=1.21.6 in
/usr/local/lib/python3.12/dist-packages (from scipy==1.11.4) (1.26.4)

```
[2]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
```

```
[3]: import scipy
scipy.__version__
```

```
[3]: '1.11.4'
```

```
[4]: from sklearn.metrics import (
    accuracy_score,
    f1_score,
    classification_report,
    confusion_matrix,
    roc_auc_score,
)
from scikitplot.metrics import plot_roc
from scikitplot.metrics import plot_precision_recall
```

```
[5]: from sklearn.datasets import load_iris

frame = load_iris(as_frame=True)
df_orig = frame['data']
X = df_orig.values
y = np.array(frame['target'])
```

```
[6]: df_orig.head()
```

```
[6]:   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
0                5.1                3.5                1.4                0.2
1                4.9                3.0                1.4                0.2
2                4.7                3.2                1.3                0.2
3                4.6                3.1                1.5                0.2
4                5.0                3.6                1.4                0.2
```

```
[7]: np.unique(y, return_counts=True)
```

```
[7]: (array([0, 1, 2]), array([50, 50, 50]))
```

```
[8]: df_orig.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sepal length (cm)      150 non-null    float64
1   sepal width (cm)       150 non-null    float64
2   petal length (cm)      150 non-null    float64
3   petal width (cm)       150 non-null    float64
dtypes: float64(4)
memory usage: 4.8 KB
```

0.1 Partitioning

```
[9]: from sklearn.model_selection import train_test_split
```

```
[10]: random_state = 0
```

```
[11]: X_train, X_test, y_train, y_test = train_test_split(
      X, y, test_size=0.4, random_state=random_state
    )
```

```
[12]: # without stratify
print(np.unique(y, return_counts=True)[1] / len(y))
print(np.unique(y_train, return_counts=True)[1] / len(y_train))
print(np.unique(y_test, return_counts=True)[1] / len(y_test))
```

```
[0.33333333 0.33333333 0.33333333]
[0.37777778 0.3         0.32222222]
[0.26666667 0.38333333 0.35      ]
```

```
[13]: X_train, X_test, y_train, y_test = train_test_split(
      X, y, test_size=0.4, stratify=y, random_state=random_state
    )
```

```
[14]: # with stratify
print(np.unique(y, return_counts=True)[1] / len(y))
print(np.unique(y_train, return_counts=True)[1] / len(y_train))
print(np.unique(y_test, return_counts=True)[1] / len(y_test))
```

```
[0.33333333 0.33333333 0.33333333]
[0.33333333 0.33333333 0.33333333]
[0.33333333 0.33333333 0.33333333]
```

```
[15]: print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

```
(90, 4) (60, 4) (90,) (60,)
```

normalization

```
[16]: from sklearn.preprocessing import StandardScaler
```

```
[17]: norm = StandardScaler()
      norm.fit(X_train)

      X_train_norm = norm.transform(X_train)
      X_test_norm = norm.transform(X_test)
```

0.2 Classifiers

Classifiers usually have a: - fit method to train them on training data - predict method to validate/test them on validation/test data

0.3 KNN

```
[18]: from sklearn.neighbors import KNeighborsClassifier
```

```
[19]: clf = KNeighborsClassifier(n_neighbors=5, metric="euclidean", weights="uniform")
      clf.fit(X_train_norm, y_train)
```

```
[19]: KNeighborsClassifier(metric='euclidean')
```

```
[20]: # predict: Predict the class labels for the provided data.
      y_test_pred = clf.predict(X_test_norm)
      y_test_pred
```

```
[20]: array([2, 2, 0, 2, 0, 0, 0, 2, 1, 1, 0, 2, 2, 0, 1, 2, 1, 0, 2, 2, 1, 0,
          2, 0, 1, 1, 2, 0, 1, 2, 0, 2, 1, 1, 2, 0, 2, 1, 1, 2, 1, 0, 1, 2,
          2, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 2, 0, 1, 1])
```

```
[21]: # y_test contains the target labels
      y_test
```

```
[21]: array([2, 2, 0, 2, 0, 0, 0, 2, 1, 1, 0, 2, 2, 0, 1, 2, 1, 0, 2, 2, 1, 0,
          2, 0, 1, 1, 2, 0, 1, 2, 0, 2, 1, 1, 2, 0, 2, 1, 1, 2, 1, 0, 1, 2,
          2, 1, 1, 0, 0, 0, 0, 2, 1, 0, 0, 1, 2, 0, 1, 1])
```

```
[22]: print("Accuracy:", accuracy_score(y_test, y_test_pred))
```

```
Accuracy: 0.9833333333333333
```

```
[23]: # score: Return the mean accuracy on the given test data and labels.
      clf.score(X_test_norm, y_test)
```

```
[23]: 0.9833333333333333
```

```
[24]: # KNeighborsClassifier.score is doing this
      (y_test_pred == y_test).sum() / len(y_test)
```

```
[24]: 0.9833333333333333
```

Performance evaluation

```
[25]: # print("F1:", f1_score(y_test, y_test_pred))
```

```
[26]: print("F1:", f1_score(y_test, y_test_pred, average="macro"))
```

```
F1: 0.9833229101521784
```

```
[27]: print("F1:", f1_score(y_test, y_test_pred, average="micro"))
```

```
F1: 0.9833333333333333
```

```
[28]: print("F1:", f1_score(y_test, y_test_pred, labels=[1], average="micro"))
```

```
F1: 0.975609756097561
```


```
[29]: print("F1:", f1_score(y_test, y_test_pred, labels=[1], average="macro"))
```

```
F1: 0.975609756097561
```

```
[30]: print(classification_report(y_test, y_test_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	20
1	0.95	1.00	0.98	20
2	1.00	0.95	0.97	20
accuracy			0.98	60
macro avg	0.98	0.98	0.98	60
weighted avg	0.98	0.98	0.98	60

```
[31]: cf = confusion_matrix(y_test, y_test_pred)
      sns.heatmap(cf, annot=True, cmap="Greens")
      plt.xlabel("True")
      plt.ylabel("Predicted")
      plt.show()
```




```
[32]: # Return probability estimates for the test data X.
y_test_pred_proba = clf.predict_proba(X_test_norm)
y_test_pred_proba[0:10]
```

```
[32]: array([[0. , 0.2, 0.8],
          [0. , 0. , 1. ],
          [1. , 0. , 0. ],
          [0. , 0. , 1. ],
          [1. , 0. , 0. ],
          [1. , 0. , 0. ],
          [1. , 0. , 0. ],
          [0. , 0. , 1. ],
          [0. , 1. , 0. ],
          [0. , 1. , 0. ]])
```

```
[33]: y_test_pred[0:10]
```

```
[33]: array([2, 2, 0, 2, 0, 0, 0, 2, 1, 1])
```


```
[34]: # https://scikit-learn.org/stable/auto\_examples/model\_selection/plot\_roc.html  
plot_roc(y_test, y_test_pred_proba)  
plt.show()
```



```
[35]: # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.  
      ↪roc\_auc\_score.html  
roc_auc_score(y_test, y_test_pred_proba, multi_class="ovr", average="macro")
```

```
[35]: 0.9987499999999999
```

```
[36]: plot_precision_recall(y_test, y_test_pred_proba)  
plt.show()
```



Repeated Holdout

```
[37]: N = 50
err = 0

for i in range(N):
    # stratified holdout
    X_rh_train, X_rh_test, y_rh_train, y_rh_test = train_test_split(X, y,
↪test_size=0.4, stratify=y)

    # normalize train set
    norm.fit(X_rh_train)
    X_rh_train_norm = norm.transform(X_rh_train)
    X_rh_test_norm = norm.transform(X_rh_test)

    # initialize and fit classifier
    clf = KNeighborsClassifier(n_neighbors=5, metric="euclidean",
↪weights="uniform")
    clf.fit(X_rh_train_norm, y_rh_train)
```



```

# computing error
acc = clf.score(X_rh_test_norm, y_rh_test)
err += 1 - acc

print("Overall error estimate:", err/N)

```

Overall error estimate: 0.04866666666666667

[37]:

Cross-validation https://scikit-learn.org/stable/modules/cross_validation.html

```

[38]: from sklearn.model_selection import cross_val_score
k = 10

```

```

[39]: # initialize classifier
clf = KNeighborsClassifier(n_neighbors=5, metric="euclidean", weights="uniform")

scores = cross_val_score(clf, X_train_norm, y_train, cv=k)
scores

```

```

[39]: array([0.88888889, 1.          , 1.          , 1.          , 0.77777778,
           1.          , 0.55555556, 0.88888889, 1.          , 1.          ])

```

```

[40]: print("Overall error estimate:", 1 - scores.mean())

```

Overall error estimate: 0.08888888888888889

```

[41]: print('Accuracy: %0.4f (+/- %0.2f)' % (scores.mean(), scores.std()))

```

Accuracy: 0.9111 (+/- 0.14)

```

[42]: # scoring default is accuracy
cross_val_score(clf, X_train_norm, y_train, cv=k, scoring='f1_macro')

```

```

[42]: array([0.88571429, 1.          , 1.          , 1.          , 0.77777778,
           1.          , 0.55555556, 0.88571429, 1.          , 1.          ])

```

[42]:

0.3.1 Hyperparameters Tuning

```

[43]: n_neighbors = range(1,20)
scores = list()

for n in n_neighbors:


```

```

    clf = KNeighborsClassifier(n_neighbors=n, metric="euclidean",
↪weights="uniform")
    clf.fit(X_train_norm, y_train)
    scores.append(clf.score(X_test_norm, y_test))

plt.plot(scores)
plt.xticks(range(len(n_neighbors)), n_neighbors)
plt.xlabel("n_neighbors")
plt.ylabel("accuracy")
plt.show()

```



```

[44]: n_neighbors = range(1,20)
      avg_scores = list()
      std_scores = list()


      for n in n_neighbors:
          clf = KNeighborsClassifier(n_neighbors=n, metric="euclidean",
↪weights="uniform")
          scores = cross_val_score(clf, X_train_norm, y_train, cv=k)
          avg_scores.append(np.mean(scores))
          std_scores.append(np.std(scores))

```

```

#plt.plot(avg_scores)
plt.errorbar(range(len(n_neighbors)), y=avg_scores, yerr=std_scores, marker='o')
plt.xticks(range(len(n_neighbors)), n_neighbors)
plt.xlabel("n_neighbors")
plt.ylabel("accuracy")
plt.show()

```



```

[45]: clf = KNeighborsClassifier(n_neighbors=12, metric="euclidean",
    ↪weights="uniform")
clf.fit(X_train_norm, y_train)
y_test_pred = clf.predict(X_test_norm)
print("Accuracy:", accuracy_score(y_test, y_test_pred))

```

Accuracy: 0.95

Grid Search

```

[46]: from sklearn.model_selection import RepeatedStratifiedKFold
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV

```

```
[47]: %%time
param_grid = {
    "n_neighbors": (np.arange(1, X_train.shape[0]//2)),
    "weights": ["uniform", "distance"],
    "metric": ["euclidean", "cityblock"],
}

grid = GridSearchCV(
    KNeighborsClassifier(),
    param_grid=param_grid,
    cv=RepeatedStratifiedKFold(random_state=0),
    n_jobs=-1,
    refit=True,
    # verbose=2
)

#param_distributions=param_grid,

grid.fit(X_train_norm, y_train)
clf = grid.best_estimator_
```

CPU times: user 2.72 s, sys: 73.8 ms, total: 2.79 s
Wall time: 41.9 s

```
[48]: print(grid.best_params_, grid.best_score_)

{'metric': 'euclidean', 'n_neighbors': 14, 'weights': 'distance'}
0.9466666666666665
```

```
[49]: y_test_pred = clf.predict(X_test_norm)
print("Accuracy:", accuracy_score(y_test, y_test_pred))
```

Accuracy: 0.9666666666666667

```
[50]: clf.score(X_test_norm, y_test)
```

```
[50]: 0.9666666666666667
```

```
[51]: # grid.cv_results_
```

```
[52]: results = pd.DataFrame(grid.cv_results_)
results
```

```
[52]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	\
0	0.004583	0.003372	0.013991	0.009098	
1	0.003988	0.002896	0.005595	0.002973	
2	0.002931	0.002279	0.009794	0.005777	
3	0.002762	0.001809	0.003857	0.001688	

4	0.002036	0.001417	0.008432	0.004344
..	...	...	...	...
171	0.001093	0.000549	0.001989	0.001017
172	0.001003	0.000344	0.001336	0.000160
173	0.001200	0.000989	0.001999	0.000790
174	0.001051	0.000410	0.001465	0.000550
175	0.001002	0.000256	0.001812	0.000439

	param_metric	param_n_neighbors	param_weights	\
0	euclidean	1	uniform	
1	euclidean	1	distance	
2	euclidean	2	uniform	
3	euclidean	2	distance	
4	euclidean	3	uniform	
..	...	...	...	
171	cityblock	42	distance	
172	cityblock	43	uniform	
173	cityblock	43	distance	
174	cityblock	44	uniform	
175	cityblock	44	distance	

	params	split0_test_score	\
0	{'metric': 'euclidean', 'n_neighbors': 1, 'wei...	0.888889	
1	{'metric': 'euclidean', 'n_neighbors': 1, 'wei...	0.888889	
2	{'metric': 'euclidean', 'n_neighbors': 2, 'wei...	0.888889	
3	{'metric': 'euclidean', 'n_neighbors': 2, 'wei...	0.888889	
4	{'metric': 'euclidean', 'n_neighbors': 3, 'wei...	0.888889	
..	...	...	
171	{'metric': 'cityblock', 'n_neighbors': 42, 'we...	0.833333	
172	{'metric': 'cityblock', 'n_neighbors': 43, 'we...	0.888889	
173	{'metric': 'cityblock', 'n_neighbors': 43, 'we...	0.833333	
174	{'metric': 'cityblock', 'n_neighbors': 44, 'we...	0.833333	
175	{'metric': 'cityblock', 'n_neighbors': 44, 'we...	0.833333	

	split1_test_score	...	split43_test_score	split44_test_score	\
0	0.944444	...	0.944444	0.833333	
1	0.944444	...	0.944444	0.833333	
2	0.888889	...	0.944444	0.888889	
3	0.944444	...	0.944444	0.833333	
4	0.888889	...	0.944444	0.888889	
..	...	...	...	...	
171	0.944444	...	0.944444	0.833333	
172	0.888889	...	0.888889	0.777778	
173	0.944444	...	0.944444	0.833333	
174	0.833333	...	0.888889	0.666667	
175	0.944444	...	0.944444	0.833333	

	split45_test_score	split46_test_score	split47_test_score	\
0	1.000000	0.944444	0.944444	
1	1.000000	0.944444	0.944444	
2	0.888889	0.944444	0.944444	
3	1.000000	0.944444	0.944444	
4	0.944444	0.944444	0.944444	
..	...	...	...	
171	0.944444	0.833333	0.944444	
172	1.000000	0.833333	0.777778	
173	1.000000	0.833333	0.944444	
174	0.888889	0.833333	0.944444	
175	0.944444	0.833333	0.944444	

	split48_test_score	split49_test_score	mean_test_score	std_test_score	\
0	0.833333	0.944444	0.911111	0.060858	
1	0.833333	0.944444	0.911111	0.060858	
2	0.833333	0.944444	0.900000	0.056656	
3	0.833333	0.944444	0.911111	0.060858	
4	0.888889	1.000000	0.925556	0.056229	
..	...	...	...	...	
171	0.944444	0.944444	0.926667	0.063285	
172	0.888889	0.888889	0.897778	0.058119	
173	0.944444	0.944444	0.930000	0.062667	
174	0.888889	0.833333	0.873333	0.069424	
175	0.944444	0.944444	0.927778	0.061111	

	rank_test_score
0	113
1	113
2	140
3	113
4	59
..	...
171	57
172	147
173	34
174	174
175	48


[176 rows x 61 columns]

```
[53]: results["metric_weight"] = results["param_metric"] + ", " +
      ↪results["param_weights"]
```

```
[54]: sns.lineplot(
      data=results, x="param_n_neighbors", y="mean_test_score",
      ↪hue="metric_weight"
```

```
)
```

```
[54]: <Axes: xlabel='param_n_neighbors', ylabel='mean_test_score'>
```



```
[54]:
```

0.3.2 Naive Bayes

```
[55]: from sklearn.naive_bayes import GaussianNB, CategoricalNB
```

Gaussian

```
[56]: clf = GaussianNB()
```

```
[57]: %%time
      clf.fit(X_train, y_train)
```

CPU times: user 2.6 ms, sys: 0 ns, total: 2.6 ms

Wall time: 2.5 ms

```
[57]: GaussianNB()
```

```
[58]: y_pred = clf.predict(X_test)
      y_pred
```

```
[58]: array([2, 2, 0, 2, 0, 0, 0, 2, 1, 1, 0, 2, 2, 0, 1, 2, 1, 0, 2, 2, 1, 0,
           2, 0, 1, 1, 2, 0, 1, 2, 0, 2, 1, 1, 2, 0, 2, 1, 1, 2, 1, 0, 1, 2,
           2, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 2, 0, 1, 2])
```

```
[59]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	20
1	0.95	0.95	0.95	20
2	0.95	0.95	0.95	20
accuracy			0.97	60
macro avg	0.97	0.97	0.97	60
weighted avg	0.97	0.97	0.97	60


```
[60]: clf.predict_proba(X_test)
```

```
[60]: array([[3.05055218e-156, 1.98469602e-002, 9.80153040e-001],
           [2.50142195e-277, 3.75932821e-012, 1.00000000e+000],
           [1.00000000e+000, 2.74048199e-017, 2.82825801e-026],
           [3.77047257e-248, 6.87918788e-010, 9.99999999e-001],
           [1.00000000e+000, 1.81924966e-016, 2.52343415e-026],
           [1.00000000e+000, 5.84024273e-017, 8.14494337e-027],
           [1.00000000e+000, 1.12344252e-016, 5.92512274e-026],
           [6.19173845e-176, 4.67941390e-004, 9.99532059e-001],
           [1.10238158e-032, 9.99999857e-001, 1.43256972e-007],
           [7.56032487e-075, 9.99968118e-001, 3.18821921e-005],
           [1.00000000e+000, 1.66546269e-010, 2.54675324e-019],
           [3.73265040e-154, 6.58305961e-002, 9.34169404e-001],
           [0.00000000e+000, 1.70481203e-012, 1.00000000e+000],
           [1.00000000e+000, 6.04531187e-012, 4.37830158e-020],
           [1.35690609e-092, 9.96128400e-001, 3.87159969e-003],
           [7.97250163e-267, 3.29344440e-010, 1.00000000e+000],
           [1.86513942e-098, 9.78842746e-001, 2.11572540e-002],
           [1.00000000e+000, 4.88720961e-019, 3.68247362e-028],
           [3.59336710e-162, 1.86623554e-002, 9.81337645e-001],
           [7.11242895e-161, 3.42090433e-002, 9.65790957e-001],
           [1.60527365e-059, 9.99978974e-001, 2.10262412e-005],
           [1.00000000e+000, 1.64906577e-013, 9.88604933e-023],
           [3.34150260e-200, 1.31058957e-006, 9.99998689e-001],
           [1.00000000e+000, 1.98789412e-017, 5.32896313e-026],
           [6.47625569e-079, 9.99895656e-001, 1.04343789e-004],
```



```
[5.29919058e-075, 9.99975461e-001, 2.45393824e-005],
[1.83093152e-230, 2.54506289e-007, 9.99999745e-001],
[1.00000000e+000, 5.44711014e-017, 2.03232679e-025],
[1.01049595e-112, 9.55075200e-001, 4.49247999e-002],
[5.37442739e-197, 6.84956056e-006, 9.99993150e-001],
[1.00000000e+000, 5.19559717e-017, 5.25698390e-026],
[1.49048596e-231, 1.04650690e-006, 9.9998953e-001],
[9.09721319e-076, 9.99853438e-001, 1.46562275e-004],
[1.18521797e-117, 7.83718534e-001, 2.16281466e-001],
[6.39795698e-170, 4.39554222e-004, 9.99560446e-001],
[1.00000000e+000, 2.48282507e-017, 6.08647478e-027],
[4.53144050e-231, 1.86017095e-008, 9.9999981e-001],
[7.95618521e-107, 9.36508854e-001, 6.34911462e-002],
[9.78193383e-052, 9.9999547e-001, 4.52961673e-007],
[1.14590516e-144, 1.41113424e-001, 8.58886576e-001],
[1.74966107e-113, 7.82539255e-001, 2.17460745e-001],
[1.00000000e+000, 4.34259424e-017, 1.51767006e-026],
[9.46439464e-108, 9.94182260e-001, 5.81774047e-003],
[7.27856849e-186, 2.70725704e-003, 9.97292743e-001],
[2.69165369e-190, 7.44082653e-004, 9.99255917e-001],
[1.93165252e-110, 8.30626172e-001, 1.69373828e-001],
[2.57768066e-078, 9.99852564e-001, 1.47435570e-004],
[1.00000000e+000, 1.44977761e-017, 1.07574272e-026],
[1.00000000e+000, 1.97922847e-016, 2.11881014e-025],
[1.00000000e+000, 8.87084357e-017, 6.08037091e-026],
[1.00000000e+000, 1.63263330e-016, 3.04083318e-026],
[4.03680708e-162, 6.49790168e-001, 3.50209832e-001],
[1.39614388e-087, 9.99941983e-001, 5.80174634e-005],
[1.00000000e+000, 3.80451594e-015, 1.43324210e-023],
[1.00000000e+000, 6.37708578e-014, 7.98165218e-022],
[1.51696008e-096, 9.99487070e-001, 5.12930491e-004],
[7.60262226e-200, 4.10479572e-006, 9.99995895e-001],
[1.00000000e+000, 1.99771588e-017, 2.03445104e-027],
[1.71249185e-067, 9.99995530e-001, 4.46979378e-006],
[4.82476701e-145, 7.05726601e-002, 9.29427340e-001]]])
```

```
[61]: # https://scikit-learn.org/stable/auto\_examples/model\_selection/plot\_roc.html
plot_roc(y_test, clf.predict_proba(X_test))
plt.show()
print(roc_auc_score(y_test, clf.predict_proba(X_test), multi_class="ovr",
↪average="macro"))
```



0.9983333333333334

0.3.3 Wine dataset

These data are the results of a chemical analysis of wines grown in the same region in Italy but derived from 3 different cultivars. The analysis determined the quantities of 13 constituents found in each of the three types of wines.

```
[62]: from sklearn.datasets import load_wine

d = load_wine(as_frame=True)
df_orig = d["data"]
y = np.array(d["target"])

df_orig.head()
```

```
[62]:   alcohol  malic_acid  ash  alcalinity_of_ash  magnesium  total_phenols  \
0    14.23         1.71  2.43              15.6         127.0           2.80
1    13.20         1.78  2.14              11.2          100.0           2.65
2    13.16         2.36  2.67              18.6          101.0           2.80
```

3	14.37	1.95	2.50	16.8	113.0	3.85
4	13.24	2.59	2.87	21.0	118.0	2.80

	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	\
0	3.06	0.28	2.29	5.64	1.04	
1	2.76	0.26	1.28	4.38	1.05	
2	3.24	0.30	2.81	5.68	1.03	
3	3.49	0.24	2.18	7.80	0.86	
4	2.69	0.39	1.82	4.32	1.04	

	od280/od315_of_diluted_wines	proline
0	3.92	1065.0
1	3.40	1050.0
2	3.17	1185.0
3	3.45	1480.0
4	2.93	735.0

```
[63]: # adding a fake categorical variable to see how to deal with it
cat_variable_1_values = ["red", "blue", "green"]
df_orig["color"] = [
    cat_variable_1_values[np.random.randint(0, len(cat_variable_1_values))]
    for _ in range(len(df_orig))
]

# creating an ordinal variable
df_orig["alkalinity_of_ash_binned"] = pd.qcut(
    df_orig["alkalinity_of_ash"], q=4, labels=False
)

X_orig = df_orig.values
df_orig.head()

# one-hot encoding of categorical data
categorical_cols = ["color"]

df = pd.get_dummies(df_orig, columns=categorical_cols)
X = df.values
df.head()
```

	alcohol	malic_acid	ash	alkalinity_of_ash	magnesium	total_phenols	\
0	14.23	1.71	2.43	15.6	127.0	2.80	
1	13.20	1.78	2.14	11.2	100.0	2.65	
2	13.16	2.36	2.67	18.6	101.0	2.80	
3	14.37	1.95	2.50	16.8	113.0	3.85	
4	13.24	2.59	2.87	21.0	118.0	2.80	

	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	\
--	------------	----------------------	-----------------	-----------------	-----	---

0	3.06	0.28	2.29	5.64	1.04
1	2.76	0.26	1.28	4.38	1.05
2	3.24	0.30	2.81	5.68	1.03
3	3.49	0.24	2.18	7.80	0.86
4	2.69	0.39	1.82	4.32	1.04

	od280/od315_of_diluted_wines	proline	alcalinity_of_ash_binned	\
0	3.92	1065.0		0
1	3.40	1050.0		0
2	3.17	1185.0		1
3	3.45	1480.0		0
4	2.93	735.0		2

	color_blue	color_green	color_red
0	False	True	False
1	False	True	False
2	False	False	True
3	False	True	False
4	False	True	False

```
[64]: X_train, X_test, y_train, y_test = train_test_split(
      X, y, test_size=0.3, stratify=y, random_state=random_state
    )
```

```
[65]: print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

```
(124, 17) (54, 17) (124,) (54,)
```

Gaussian NB

```
[66]: clf = GaussianNB()
```

```
[67]: %%time
      clf.fit(X_train, y_train)
```

```
CPU times: user 2.28 ms, sys: 758 µs, total: 3.04 ms
```

```
Wall time: 3.03 ms
```

```
[67]: GaussianNB()
```


```
[68]: y_pred = clf.predict(X_test)
```

```
[69]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.94	0.97	18
1	0.95	0.95	0.95	21

	2	0.94	1.00	0.97	15
accuracy				0.96	54
macro avg		0.96	0.97	0.96	54
weighted avg		0.96	0.96	0.96	54

```
[70]: # https://scikit-learn.org/stable/auto\_examples/model\_selection/plot\_roc.html
plot_roc(y_test, clf.predict_proba(X_test))
plt.show()
print(roc_auc_score(y_test, clf.predict_proba(X_test), multi_class="ovr",
↪average="macro"))
```



0.9995189995189996

[70]:

Categorical

```

[71]: non_cat_columns = [
        "alcohol",
        "malic_acid",
        "ash",
        "alcalinity_of_ash",
        "magnesium",
        "total_phenols",
        "flavanoids",
        "nonflavanoid_phenols",
        "proanthocyanins",
        "color_intensity",
        "hue",
        "od280/od315_of_diluted_wines",
        "proline",
    ]

X_noncat = df[non_cat_columns].values

X_train_noncat, X_test_noncat, y_train_noncat, y_test_noncat = train_test_split(
    X_noncat, y, test_size=0.3, stratify=y, random_state=random_state
)

# train and test set should be binned separately
X_train_cat = list()
for column_idx in range(X_train_noncat.shape[1]):
    X_train_cat.append(pd.qcut(X_train_noncat[:, column_idx], q=4,
        labels=False))
X_train_cat = np.array(X_train_cat).T

X_test_cat = list()
for column_idx in range(X_test_noncat.shape[1]):
    X_test_cat.append(pd.qcut(X_test_noncat[:, column_idx], q=4, labels=False))
X_test_cat = np.array(X_test_cat).T

print(X_train_cat.shape, X_test_cat.shape)

```

(124, 13) (54, 13)

```
[72]: X_train_cat
```

```

[72]: array([[2, 3, 1, ..., 1, 1, 1],
        [3, 0, 3, ..., 3, 2, 3],
        [3, 0, 3, ..., 3, 1, 3],
        ...,
        [3, 3, 0, ..., 0, 0, 0],
        [0, 0, 1, ..., 2, 2, 2],
        [1, 3, 2, ..., 3, 2, 0]])

```

```
[73]: clf = CategoricalNB()
      clf.fit(X_train_cat, y_train_noncat)
```

```
[73]: CategoricalNB()
```

```
[74]: X_train_cat
```

```
[74]: array([[2, 3, 1, ..., 1, 1, 1],
           [3, 0, 3, ..., 3, 2, 3],
           [3, 0, 3, ..., 3, 1, 3],
           ...,
           [3, 3, 0, ..., 0, 0, 0],
           [0, 0, 1, ..., 2, 2, 2],
           [1, 3, 2, ..., 3, 2, 0]])
```

```
[75]: y_pred = clf.predict(X_test_cat)
```

```
[76]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	18
1	1.00	0.90	0.95	21
2	0.88	1.00	0.94	15
accuracy			0.96	54
macro avg	0.96	0.97	0.96	54
weighted avg	0.97	0.96	0.96	54

1 Binary Classification Example

```
[77]: from scikitplot.metrics import plot_cumulative_gain, plot_lift_curve
      # plot_cumulative_gain and plot_lift_curve only work in a binary classification_
      ↪ case
```

```
[78]: from sklearn.datasets import load_breast_cancer
```

```
[79]: X, y = load_breast_cancer(return_X_y=True)
```

```
[79]:
```

```
[80]: X_train, X_test, y_train, y_test = train_test_split(
      X, y, test_size=0.3, stratify=y, random_state=random_state
      )
```

```
[81]: clf = GaussianNB()
```

```
[82]: clf.fit(X_train, y_train)
```

```
[82]: GaussianNB()
```

```
[83]: y_test_pred_proba = clf.predict_proba(X_test)
      y_test_pred_proba
```


```
[83]: array([[1.00000000e+000, 9.60798304e-067],
             [3.28452382e-012, 1.00000000e+000],
             [1.00000000e+000, 1.37358518e-029],
             [9.07175634e-001, 9.28243663e-002],
             [1.00000000e+000, 1.13066573e-026],
             [4.86470493e-011, 1.00000000e+000],
             [1.00000000e+000, 1.85354720e-229],
             [1.00000000e+000, 2.24061507e-034],
             [1.00000000e+000, 7.64991022e-020],
             [1.00000000e+000, 5.98723466e-025],
             [1.55676568e-010, 1.00000000e+000],
             [4.87665583e-012, 1.00000000e+000],
             [2.51652415e-014, 1.00000000e+000],
             [3.35003877e-015, 1.00000000e+000],
             [1.00000000e+000, 7.27215795e-019],
             [1.00000000e+000, 5.01138848e-047],
             [1.00000000e+000, 3.84869046e-038],
             [5.11581566e-005, 9.99948842e-001],
             [1.00000000e+000, 5.99928618e-030],
             [1.00000000e+000, 8.93234543e-263],
             [9.23899517e-016, 1.00000000e+000],
             [1.22947094e-013, 1.00000000e+000],
             [4.60035794e-013, 1.00000000e+000],
             [1.00000000e+000, 2.83284900e-302],
             [1.00000000e+000, 1.92821921e-054],
             [4.94411999e-005, 9.99950559e-001],
             [6.63793002e-012, 1.00000000e+000],
             [1.00000000e+000, 4.93986527e-097],
             [5.84717888e-014, 1.00000000e+000],
             [1.00000000e+000, 1.59716799e-098],
             [8.87809741e-001, 1.12190259e-001],
             [4.04806236e-013, 1.00000000e+000],
             [9.99999683e-001, 3.16530138e-007],
             [3.93236381e-014, 1.00000000e+000],
             [1.25847936e-014, 1.00000000e+000],
             [1.00000000e+000, 1.61535881e-014],
             [6.30346164e-008, 9.99999937e-001],
             [1.00000000e+000, 9.79137325e-237],
```


[3.98832803e-019, 1.00000000e+000],
[9.96202524e-001, 3.79747602e-003],
[6.37533647e-001, 3.62466353e-001],
[1.00000000e+000, 3.47386136e-073],
[8.65947829e-011, 1.00000000e+000],
[1.00000000e+000, 2.16768230e-012],
[4.79603200e-010, 1.00000000e+000],
[1.00000000e+000, 8.70229084e-065],
[1.00000000e+000, 1.49902651e-028],
[2.13772341e-016, 1.00000000e+000],
[2.42737760e-019, 1.00000000e+000],
[1.27400030e-016, 1.00000000e+000],
[4.08428604e-010, 1.00000000e+000],
[1.00000000e+000, 3.06395512e-201],
[6.37932236e-018, 1.00000000e+000],
[1.00000000e+000, 8.43935225e-045],
[2.87486870e-005, 9.99971251e-001],
[1.68740960e-015, 1.00000000e+000],
[9.79395040e-001, 2.06049605e-002],
[6.30316863e-008, 9.9999937e-001],
[1.95612618e-014, 1.00000000e+000],
[2.47412098e-014, 1.00000000e+000],
[6.41925512e-007, 9.9999358e-001],
[8.87706008e-015, 1.00000000e+000],
[1.00000000e+000, 7.27668630e-020],
[1.00000000e+000, 5.75617721e-022],
[1.00000000e+000, 4.83161832e-017],
[9.99571463e-001, 4.28537207e-004],
[1.00000000e+000, 2.76760752e-014],
[1.00000000e+000, 5.84136695e-100],
[1.00000000e+000, 4.19305275e-026],
[9.96580698e-019, 1.00000000e+000],
[7.84186796e-020, 1.00000000e+000],
[1.00000000e+000, 1.62440933e-055],
[5.11529722e-017, 1.00000000e+000],
[6.10023118e-010, 9.9999999e-001],
[3.53918511e-010, 1.00000000e+000],
[3.89701907e-014, 1.00000000e+000],
[2.76850277e-010, 1.00000000e+000],
[9.99999984e-001, 1.58812530e-008],
[1.00000000e+000, 6.62179158e-015],
[6.07278434e-009, 9.9999994e-001],
[1.86299027e-013, 1.00000000e+000],
[2.57800545e-015, 1.00000000e+000],
[8.67048632e-009, 9.9999991e-001],
[1.00000000e+000, 8.20268595e-021],
[3.06330200e-001, 6.93669800e-001],


[2.76904635e-012, 1.00000000e+000],
[1.00000000e+000, 1.07090316e-030],
[1.00000000e+000, 5.88202001e-025],
[1.00000000e+000, 1.12952567e-049],
[1.00000000e+000, 1.37101315e-090],
[1.00000000e+000, 5.06841370e-012],
[6.87390643e-004, 9.99312609e-001],
[2.32024148e-014, 1.00000000e+000],
[7.98351890e-007, 9.99999202e-001],
[2.45345747e-010, 1.00000000e+000],
[1.00000000e+000, 4.96544130e-047],
[1.00000000e+000, 2.73594696e-026],
[3.34701403e-008, 9.9999967e-001],
[1.00000000e+000, 2.08754217e-057],
[9.9999998e-001, 2.43932083e-009],
[9.64990534e-003, 9.90350095e-001],
[8.02821659e-014, 1.00000000e+000],
[2.26325644e-011, 1.00000000e+000],
[1.00000000e+000, 2.69374767e-011],
[7.90120703e-012, 1.00000000e+000],
[1.13204641e-002, 9.88679536e-001],
[4.93482495e-011, 1.00000000e+000],
[9.99999878e-001, 1.22307191e-007],
[4.50662369e-012, 1.00000000e+000],
[8.25753085e-011, 1.00000000e+000],
[3.07725427e-015, 1.00000000e+000],
[1.00000000e+000, 2.85640504e-044],
[2.51559739e-002, 9.74844026e-001],
[4.61383390e-011, 1.00000000e+000],
[1.82186521e-016, 1.00000000e+000],
[5.28308302e-006, 9.99994717e-001],
[1.65228830e-011, 1.00000000e+000],
[1.00000000e+000, 4.56498688e-022],
[1.23132066e-017, 1.00000000e+000],
[1.22800876e-015, 1.00000000e+000],
[3.56992997e-014, 1.00000000e+000],
[2.03488382e-013, 1.00000000e+000],
[2.81915044e-016, 1.00000000e+000],
[2.36902189e-019, 1.00000000e+000],
[9.99995090e-001, 4.91045622e-006],
[4.83312761e-013, 1.00000000e+000],
[2.31509954e-014, 1.00000000e+000],
[9.32648265e-014, 1.00000000e+000],
[7.72754547e-016, 1.00000000e+000],
[8.43325608e-015, 1.00000000e+000],
[1.25510491e-017, 1.00000000e+000],
[1.29458768e-012, 1.00000000e+000],

```
[6.16172274e-016, 1.00000000e+000],
[5.64369777e-012, 1.00000000e+000],
[1.00000000e+000, 2.20873654e-061],
[7.71743562e-014, 1.00000000e+000],
[2.00824407e-011, 1.00000000e+000],
[2.77403674e-001, 7.22596326e-001],
[1.00000000e+000, 2.19534740e-044],
[1.00000000e+000, 2.50277461e-043],
[6.29379257e-018, 1.00000000e+000],
[1.00000000e+000, 2.98256222e-110],
[1.06730721e-019, 1.00000000e+000],
[1.00000000e+000, 2.50446769e-112],
[6.39069091e-013, 1.00000000e+000],
[1.00000000e+000, 1.47862500e-054],
[5.65173547e-018, 1.00000000e+000],
[8.37303002e-016, 1.00000000e+000],
[5.84731461e-015, 1.00000000e+000],
[2.86749633e-015, 1.00000000e+000],
[7.62722007e-016, 1.00000000e+000],
[7.16277436e-015, 1.00000000e+000],
[6.90868938e-014, 1.00000000e+000],
[5.36775578e-006, 9.99994632e-001],
[1.00000000e+000, 1.09686447e-082],
[2.21035391e-015, 1.00000000e+000],
[5.08074190e-004, 9.99491926e-001],
[8.51190961e-018, 1.00000000e+000],
[1.30925836e-016, 1.00000000e+000],
[1.24317330e-009, 9.99999999e-001],
[1.00000000e+000, 2.47391080e-111],
[1.87794810e-016, 1.00000000e+000],
[2.36997051e-011, 1.00000000e+000],
[8.62804098e-014, 1.00000000e+000],
[5.10489927e-003, 9.94895101e-001],
[2.25348614e-017, 1.00000000e+000],
[5.95648811e-011, 1.00000000e+000],
[7.42836733e-003, 9.92571633e-001],
[2.03031410e-011, 1.00000000e+000],
[2.39826926e-012, 1.00000000e+000],
[1.02708885e-022, 1.00000000e+000]]])
```

```
[84]: plot_roc(y_test,y_test_pred_proba)
plt.show()
```




```
[85]: plot_precision_recall(y_test, y_test_pred_proba)
plt.show()
```



For binary classification tasks you can also try to use cumulative gains curve and lift curve as implemented here

```
[86]: fig, axs = plt.subplots(2, 1, sharex=True, figsize=(8,6))
      plot_cumulative_gain(y_test, y_test_pred_proba, ax=axs[0])
      plot_lift_curve(y_test, y_test_pred_proba, ax=axs[1])
      plt.tight_layout()
      plt.show()
```



[86] :

[86] :